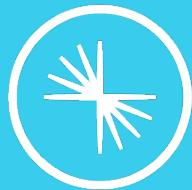


Apache Kafka® Administration by Confluent

Exercise Book

Version 6.0.0-v2.0.1



CONFLUENT

Table of Contents

Copyright & Trademarks	1
Lab 01 Fundamentals of Apache Kafka	2
a. Introduction	2
b. Using Kafka's Command-Line Tools	15
c. Producing Records with a Null Key	21
Lab 04 Providing Durability	26
a. Investigating the Distributed Log	26
Lab 05 Configuring a Kafka Cluster	46
a. Exploring Configuration	46
b. Increasing Replication Factor	51
Lab 06 Managing a Kafka Cluster	55
a. Kafka Administrative Tools	55
Lab 07 Consumer Groups and Load Balancing	68
a. Modifying Partitions and Viewing Offsets	68
Lab 08 Optimizing Kafka's Performance	74
a. Exploring Producer Performance	74
b. Performance Tuning	76
Lab 09 Kafka Security	92
a. Securing the Kafka Cluster	92
Lab 10 Data Pipelines with Kafka Connect	103
a. Running Kafka Connect	103
Appendix A: Reassigning Partitions in a Topic - Alternate Method	115
Appendix B: Running Labs in Docker for Desktop	121

Copyright & Trademarks

Copyright © Confluent, Inc. 2014-2022. [Privacy Policy](#) | [Terms & Conditions](#).

Apache, Apache Kafka, Kafka and the Kafka logo are trademarks of the

[Apache Software Foundation](#)

Lab 01 Fundamentals of Apache Kafka

a. Introduction

This document provides Hands-On Exercises for the course **Apache Kafka® Administration by Confluent**. You will use a setup that includes a virtual machine (VM) configured as a Docker host to demonstrate the distributed nature of Apache Kafka.

The main Kafka cluster includes the following components, each running in a Docker container:

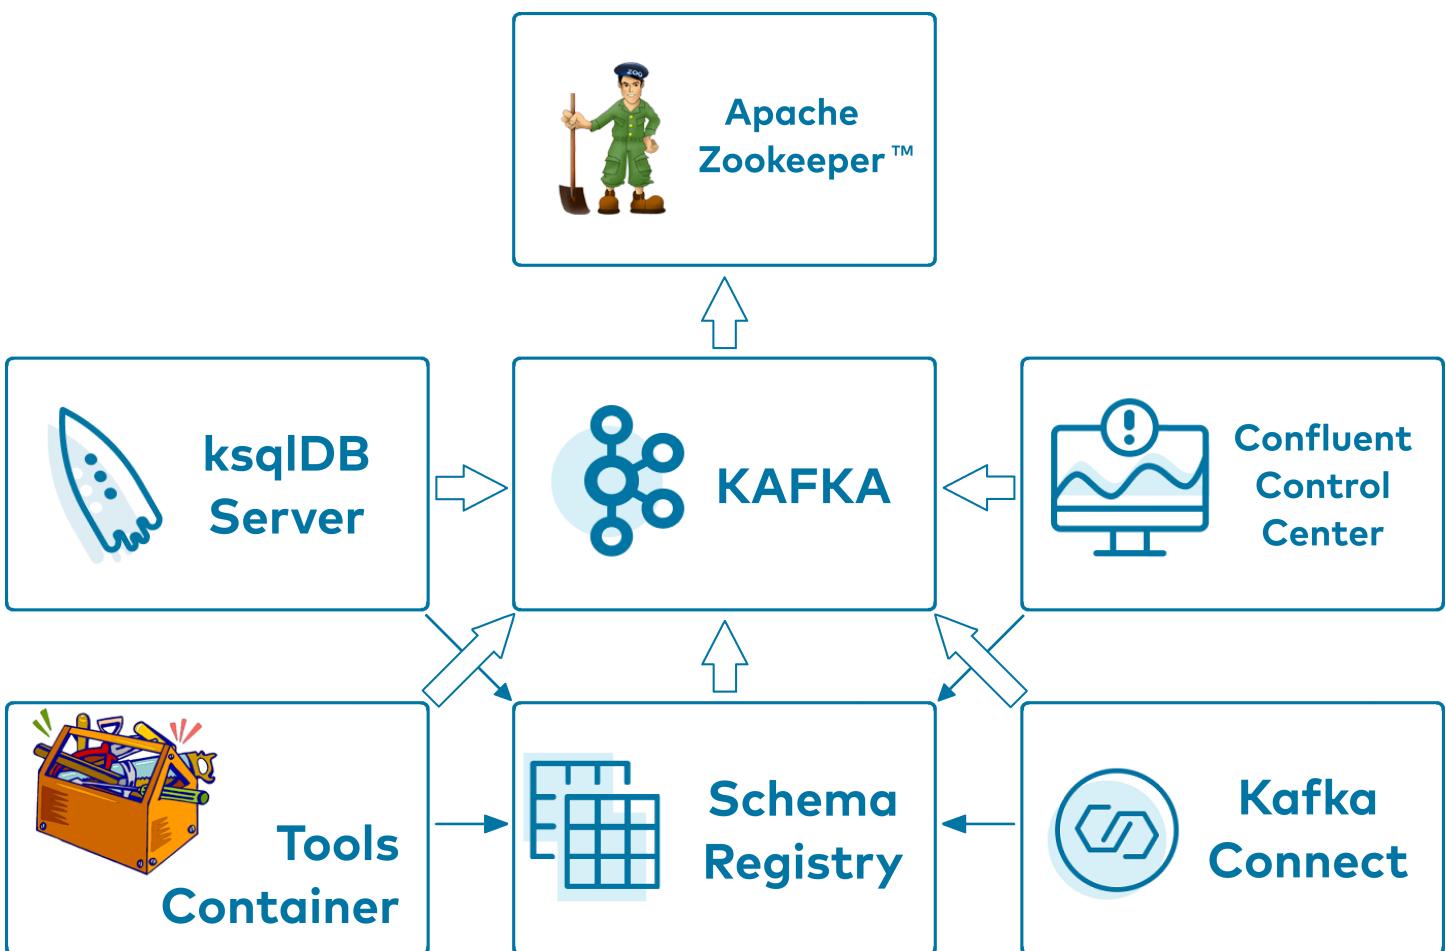


Table 1. Components of the Confluent Platform

Hostname	Description
zk-1	ZooKeeper 1
zk-2	ZooKeeper 2

Hostname	Description
zk-3	ZooKeeper 3
kafka-1	Kafka Broker 1
kafka-2	Kafka Broker 2
kafka-3	Kafka Broker 3
schema-registry	Schema Registry
connect	Kafka Connect
ksqlDB-server	ksqlDB Server
control-center	Confluent Control Center
tools	A container used to run tools against the cluster

As you progress though the exercises you will selectively turn on parts of your cluster as they are needed.

You will use Confluent Control Center to monitor the main Kafka cluster. To achieve this, we are also running the Control Center service which is backed by the same Kafka cluster.

In this course we are using Confluent Platform version 6.0.0 which includes Kafka 2.6.0.



In production, Control Center should be deployed with its own dedicated Kafka cluster, separate from the cluster with production traffic. Using a dedicated metrics cluster is more resilient because it continues to provide system health monitoring even if the production traffic cluster experiences issues.

Alternative Lab Environments

As an alternative you can also download the VM to your laptop and run it in VirtualBox. Make sure you have the newest version of VirtualBox installed. Download the VM from this link:

- <https://confluent-training-images-us-east-1.s3.amazonaws.com/training-ubuntu-18-04-may2020.ova>

If you have installed Docker for Desktop on your Mac or Windows 10 Pro machine, then you can run the labs there. But please note that your trainer might not be able to troubleshoot any potential problems if you are running the labs locally. If you choose to do this, follow the instructions at → [Running Labs in Docker for Desktop](#).

Command Line Examples

Most Exercises contain commands that must be run from the command line. These commands will look like this:

```
$ pwd  
/home/training
```

Commands you should type are shown in bold; non-bold text is an example of the output produced as a result of the command.

Continued Learning After Class

Once the course ends, the VM in Content Raven will terminate and you will no longer have access to it. However, you can still download the VM onto your own machine or use Docker locally to revisit these materials. We encourage you to bring up your own test environment, explore configuration files, inspect scripts, and perform tests. Here are some activities we encourage to reinforce your learning:

- Revisit the exercises in this exercise book
- Summarize and discuss the student handbook with your peers
- Consult the README in this course's public source repository:
<https://github.com/confluentinc/training-operations-src>

Preparing the Lab

Welcome to your lab environment! You are connected as user **training**, password **training**.

If you haven't already done so, you should open the **Exercise Guide** that is located on the lab virtual machine. To do so, open the **Confluent Training Exercises** folder that is located on the lab virtual machine desktop. Then double-click the shortcut that is in the folder to open the **Exercise Guide**.



Copy and paste works best if you copy from the Exercise Guide on your lab virtual machine.

- Standard Ubuntu keyboard shortcuts will work: **Ctrl+C** → Copy, **Ctrl+V** → Paste
- In a Terminal window: **Ctrl+Shift+C** → Copy, **Ctrl+Shift+V** → Paste.

If you find these keyboard shortcuts are not working you can use the right-click context menu for copy and paste.

1. Open a terminal window
2. Clone the source code repository to the folder **confluent-admin** in your **home** directory and change to that directory:

```
$ cd ~  
$ git clone --branch 6.0.0-v2.0.1 \  
https://github.com/confluentinc/training-administration-src.git \  
confluent-admin
```



If you chose to select another folder for the labs then note that many of our samples assume that the lab folder is [~/confluent-admin](#). You will have to adjust all those commands to fit your specific environment.

3. Navigate to the **confluent-admin** folder:

```
$ cd ~/confluent-admin
```

4. Start the complete Kafka cluster with the following command:

```
$ docker-compose up -d  
Creating network "confluent-admin_confluent" with the default driver  
Creating volume "confluent-admin_data-zk-log-1" with default driver  
Creating volume "confluent-admin_data-zk-data-1" with default driver  
Creating volume "confluent-admin_data-zk-log-2" with default driver  
Creating volume "confluent-admin_data-zk-data-2" with default driver  
Creating volume "confluent-admin_data-zk-log-3" with default driver  
Creating volume "confluent-admin_data-zk-data-3" with default driver  
Creating volume "confluent-admin_data-kafka-1" with default driver  
Creating volume "confluent-admin_data-kafka-2" with default driver  
Creating volume "confluent-admin_data-kafka-3" with default driver  
Creating zk-2 ... done  
Creating schema-registry ... done  
Creating connect ... done  
Creating tools ... done  
Creating kafka-2 ... done  
Creating kafka-1 ... done  
Creating ksqldb-server ... done  
Creating zk-1 ... done  
Creating zk-3 ... done  
Creating kafka-3 ... done  
Creating control-center ... done
```

5. Monitor the cluster with:

Name	Command	State
Ports		
connect	/etc/confluent/docker/run	Up (healthy)
0.0.0.0:8083->8083/tcp, 9092/tcp		
control-center	/etc/confluent/docker/run	Up
0.0.0.0:9021->9021/tcp		
kafka-1	/etc/confluent/docker/run	Up
0.0.0.0:19092->19092/tcp, 9092/tcp		
kafka-2	/etc/confluent/docker/run	Up
0.0.0.0:29092->29092/tcp, 9092/tcp		
kafka-3	/etc/confluent/docker/run	Up
0.0.0.0:39092->39092/tcp, 9092/tcp		
ksqldb-server	/etc/confluent/docker/run	Up (healthy)
0.0.0.0:8088->8088/tcp		
schema-registry	/etc/confluent/docker/run	Up
0.0.0.0:8081->8081/tcp		
tools	/bin/bash	Up
zk-1	/etc/confluent/docker/run	Up
0.0.0.0:12181->12181/tcp, 2181/tcp, 2888/tcp, 3888/tcp		
zk-2	/etc/confluent/docker/run	Up
2181/tcp, 0.0.0.0:22181->22181/tcp, 2888/tcp, 3888/tcp		
zk-3	/etc/confluent/docker/run	Up
2181/tcp, 2888/tcp, 0.0.0.0:32181->32181/tcp, 3888/tcp		

All services should have **State** equal to **Up**.

6. OPTIONAL: You can also observe the stats of Docker on your VM:

```
$ docker stats
```

CONTAINER ID / LIMIT	NAME	CPU %	MEM USAGE
MEM %	...		
23a0817edf17 / 9.758GiB	kafka-3	51.45%	450.4MiB
2133c07a5909 / 9.758GiB	control-center	42.19%	443.5MiB
632188f5590d / 9.758GiB	zk-3	0.09%	88.41MiB
ad213dd14b53 / 9.758GiB	ksqlDB-server	0.56%	252.8MiB
71709f3da755 / 9.758GiB	zk-1	0.09%	88.86MiB
da2fe67720c1 / 9.758GiB	kafka-1	6.27%	458.8MiB
b9852c134792 / 9.758GiB	tools	0.00%	1.273MiB
842fa80a9c2f / 9.758GiB	kafka-2	35.25%	446.7MiB
c04a39218779 / 9.758GiB	connect	0.80%	1.629GiB
fe51891d9ae5 / 9.758GiB	schema-registry	0.31%	180.4MiB
	zk-2	0.09%	92.99MiB
	...		

Press Ctrl+C to exit Docker stats.

Testing the Installation

1. Use the **zookeeper-shell** command to verify that all Brokers have registered with ZooKeeper. You should see the three Brokers listed as **[101, 102, 103]** in the last line of the output.

```
$ zookeeper-shell zk-1:12181 ls /brokers/ids
Connecting to zk-1:12181

WATCHER::  
WatchedEvent state:SyncConnected type:None path:null  
[101, 102, 103]
```

 For those **not** using the VM but instead are using Docker for Desktop, remember to run commands against the cluster from within the **tools** container with **docker-compose exec tools bash**.

OPTIONAL: Analyzing the Docker Compose File

1. Open the file **docker-compose.yml** in your editor and:
 - a. locate the various services that are listed in the table earlier in this section
 - b. note that the **hostname** (e.g. **zk-1** or **kafka-2**) is used to resolve a particular service
 - c. note how each ZooKeeper instance (zk-1, zk-2, zk-3)
 - i. gets a unique ID assigned via environment variable **ZOOKEEPER_SERVER_ID**
 - ii. gets the information about all members of the ensemble:

```
ZOOKEEPER_SERVERS=zk-1:2888:3888;zk-2:2888:3888;zk-3:2888:3888
```

- d. note how each Broker (kafka-1, kafka-2, kafka-3)
 - i. gets a unique ID assigned via environment variable **KAFKA_BROKER_ID**
 - ii. defines where to find the ZooKeeper instances

```
KAFKA_ZOOKEEPER_CONNECT: zk-1:12181,zk-2:22181,zk-3:32181
```

- iii. configures the Broker to send metrics to Control Center:

```
KAFKA_METRIC_REPORTERS:  
"io.confluent.metrics.reporter.ConfluentMetricsReporter"  
CONFLUENT_METRICS_REPORTER_BOOTSTRAP_SERVERS: "kafka-  
1:9092,kafka-2:9092,kafka-3:9092"
```

- e. note how various services use the environment variable **..._BOOTSTRAP_SERVERS** to define the list of Kafka Brokers that serve as bootstrap servers:

```
..._BOOTSTRAP_SERVERS: kafka-1:9092,kafka-2:9092,kafka-3:9092
```

- f. note how e.g. the **connect** service and the **ksql-server** service define Producer and Consumer interceptors that produce data which can be monitored in Control Center:

```
io.confluent.monitoring.clients.interceptor.MonitoringProducerInte  
rceptor  
io.confluent.monitoring.clients.interceptor.MonitoringConsumerInte  
rceptor
```

Using Control Center

1. In the VM, open a new browser tab in Google Chrome.
2. Navigate to Control Center at <http://localhost:9021>:

The screenshot shows the Confluent Control Center Home page. At the top, there's a header with the Confluent logo, a trial notice ("Enterprise trial ends in 28 days"), and a bell icon. Below the header, the word "Home" is displayed in large letters. To the left, there's a sidebar with icons for "Clusters" and "Control Center". The main area shows a summary: "1 Healthy clusters" and "0 Unhealthy clusters". There's also a search bar labeled "Search cluster name" and a toggle switch labeled "Hide healthy clusters". A detailed cluster card is expanded for "controlcenter.cluster", which is listed as "Running". The card contains an "Overview" section with metrics: Brokers (3), Partitions (246), Topics (44), Production (33.5kB/s), and Consumption (26.7kB/s). It also lists "Connected services" with KSQL clusters (1) and Connect clusters (1).



The Control Center environment will take a few minutes to stabilize and report a healthy cluster.

3. Select the cluster **Cluster 1** and you will see this:

Overview

Brokers

3	32.1k	26.4k
Total	Production (bytes / second)	Consumption (bytes / second)

Topics

43	245	0	0
Total	Partitions	Under replicated partitions	Out-of-sync replicas

Connect

1	0	0	0	0
Clusters	Running	Paused	Degraded	Failed

4. Click on **Brokers** on the left to go to the **Brokers overview** page.

Brokers overview

Production

32.7k bytes / second

Consumption

27.1k bytes / second

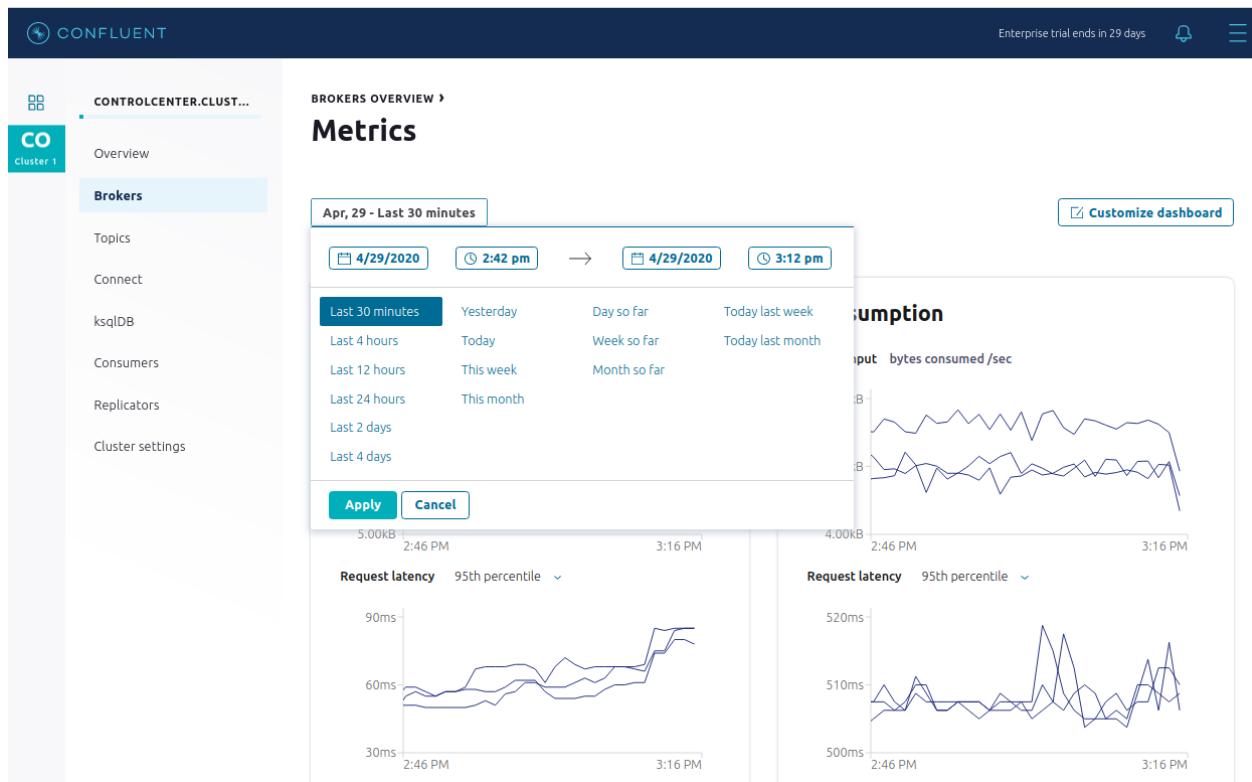
Partitioning and replication

45	253	747
Topics	Partitions	Replicas

Active controller
broker.id 101
Broker ID

ZooKeeper
Yes
ZooKeeper connected

5. In the **Brokers overview** page, click on **Production** to open the **Metrics** page.
6. Click on the button that shows the current date, and change **Last 4 hours** to **Last 30 minutes**.



- In the **Production** section of the **Metrics** page, you are looking at metrics for **Produce** requests in your cluster. Hover your mouse over the **Throughput** and **Request latency** graphs.

Consumption metrics are visible either to the right of or below the **Production** metrics depending upon the Control Center window width. This section of the **Metrics** page displays similar metrics for **Consume** requests.



In production, Control Center should be deployed with its own dedicated Kafka cluster, separate from the cluster with production traffic. Using a dedicated metrics cluster is more resilient because it continues to provide system health monitoring even if the production traffic cluster experiences issues.

Cleanup

- Execute the following command to completely cleanup your environment:

```
$ docker-compose down -v
```

b. Using Kafka's Command-Line Tools

In this Hands-On Exercise you will start to become familiar with some of Kafka's command-line tools. Specifically you will:

- Use a tool to **create** a Topic
- Use a console program to **produce** a message
- Use a console program to **consume** a message
- Use a tool to explore data stored in ZooKeeper

Prerequisites

1. Navigate to the **confluent-admin** folder:

```
$ cd ~/confluent-admin
```

2. Run the Kafka cluster, including Control Center:

```
$ docker-compose up -d
```

Console Producing and Consuming

Kafka has built-in command line utilities to produce messages to a Topic and read messages from a Topic. These are extremely useful to verify that Kafka is working correctly, and for testing and debugging.

1. Before we can start writing data to a Topic in Kafka, we need to first create that Topic using a tool called **kafka-topics**. From within the terminal window run the command:

```
$ kafka-topics
```

This will bring up a list of parameters that the **kafka-topics** program can receive. Take a moment to look through the options.

- Now execute the following command to create the Topic **testing**:

```
$ kafka-topics --bootstrap-server kafka-1:19092 \
  --create \
  --partitions 1 \
  --replication-factor 1 \
  --topic testing
```

We create the Topic with a single Partition and **replication-factor** of one.



If **auto.create.topics.enable** was set to **true** and the previous step had not been completed, the Topic would automatically be created when the first record is written to it using the default topic number of partitions and default replication factor. Enabling auto topic creation is **strongly discouraged** in production. Always create your Topics explicitly!

- Now let's move on to start writing data into the Topic just created. From within the terminal window run the command:

```
$ kafka-console-producer
```

This will bring up a list of parameters that the **kafka-console-producer** program can receive. Take a moment to look through the options. We will discuss many of their meanings later in the course.

- Run **kafka-console-producer** again with the required arguments:

```
$ kafka-console-producer --bootstrap-server kafka-1:19092 --topic
testing
```

The tool prompts you with a **>**.

- At this prompt type:

```
> some data
```

And hit **Enter**.

6. Now type:

```
> more data
```

And hit **Enter**.

7. Type:

```
> final data
```

And hit **Enter**.

8. Press **Ctrl+D** to exit the **kafka-console-producer** program.

9. Now we will use a Consumer to retrieve the data that was produced. Run the command:

```
$ kafka-console-consumer
```

This will bring up a list of parameters that the **kafka-console-consumer** can receive. Take a moment to look through the options.

10. Run **kafka-console-consumer** again with the following arguments:

```
$ kafka-console-consumer \
  --bootstrap-server kafka-1:19092 \
  --from-beginning \
  --topic testing
```

After a short moment you should see all the messages that you produced using **kafka-console-producer** earlier:

```
some data
more data
final data
```

11. Press **Ctrl+C** to exit **kafka-console-consumer**.

OPTIONAL: Running Producer and Consumer in Parallel

The **kafka-console-producer** and **kafka-console-consumer** programs can be run at the same time. Run **kafka-console-producer** and **kafka-console-consumer** in separate terminal windows at the same time to see how **kafka-console-consumer** receives the events.

OPTIONAL: Working with record keys

By default, **kafka-console-producer** and **kafka-console-consumer** assume null keys. They can also be run with appropriate arguments to write and read keys as well as values.

1. Re-run the Producer with additional arguments to write (key,value) pairs to the Topic:

```
$ kafka-console-producer \
  --bootstrap-server kafka-1:19092 \
  --topic testing \
  --property parse.key=true \
  --property key.separator=,
```

2. Enter a few values such as:

```
> 1,my first record
> 2,another record
> 3,Kafka is cool
```

3. Press **Ctrl+D** to exit the producer.

4. Now run the **Consumer** with additional arguments to print the key as well as the value:

```
$ kafka-console-consumer \
  --bootstrap-server kafka-1:19092 \
  --from-beginning \
  --topic testing \
  --property print.key=true

null      some data
null      more data
null      final data
1        my first record
2        another record
3        Kafka is cool
```

Note the **NULL** values for the first 3 records that we entered earlier...

5. Press **Ctrl+C** to exit the Consumer.

The ZooKeeper Shell

1. Access Kafka's data in ZooKeeper by using the **zookeeper-shell** command to connect to the ZooKeeper host **zk-1**:

```
$ zookeeper-shell zk-1:12181
Connecting to zookeeper
Welcome to ZooKeeper!
JLine support is disabled

WATCHER:::

WatchedEvent state:SyncConnected type:None path:null
```

2. From within the **zookeeper-shell** application, type **ls /** to view the directory structure in ZooKeeper. Note the **/** is required.

```
ls /
[admin, brokers, cluster, config, consumers, controller,
controller_epoch, isr_change_notification, latest_producer_id_block,
log_dir_event_notification, zookeeper]
```

3. Type **ls /brokers** to see this next level of the directory structure.

```
ls /brokers
[ids, seqid, topics]
```

4. Type **ls /brokers/ids** to see the broker ids for the Kafka cluster.

```
ls /brokers/ids
[101,102,103]
```

Note the last output **[101, 102, 103]**, indicating that we have 3 Brokers with IDs **101, 102, 103** in our cluster.

5. Type **get /brokers/ids/101** to see the metadata for broker 101.

```
get /brokers/ids/101
{"listener_security_protocol_map": {"DOCKER": "PLAINTEXT", "HOST": "PLAINTEXT"}, "endpoints": ["DOCKER://kafka-1:9092", "HOST://kafka-1:9092"], "jmx_port": -1, "port": 9092, "host": "kafka-1", "version": 4, "timestamp": "1602909525108"}
```

6. Type **get /brokers/topics/testing/partitions/0/state** to see the metadata for partition 0 of topic **testing**.

```
get /brokers/topics/testing/partitions/0/state
{"controller_epoch": 1, "leader": 103, "version": 1, "leader_epoch": 0, "isr": [103]}
```

Note: During client startup, it requests cluster metadata from a broker in the **bootstrap.servers** list. The output of the two previous commands reflects a bit of this cluster metadata included in the broker response.

7. Press **Ctrl+D** to exit the ZooKeeper shell.

Cleanup

1. Execute the following command to completely cleanup your environment:

```
$ docker-compose down -v
```

Conclusion

In this lab you have used Kafka command line tools to create a Topic, write and read from this Topic. Finally you have used the ZooKeeper shell tool to access data stored within ZooKeeper.

c. Producing Records with a Null Key

In this Hands-On Exercise, you will create a Topic with **multiple Partitions**, produce records with a **null** key to those Partitions using the default partitioner, and then read it back to observe issues with ordering.

Prerequisites

1. Please make sure you have prepared your environment by following → [Preparing the Labs](#)
2. Start the Kafka cluster:

```
$ cd ~/confluent-admin  
$ docker-compose up -d
```

Produce to multiple Partitions

1. From within the terminal window create a Topic manually with Kafka's command-line tool, specifying that it should have **two Partitions**:

```
$ kafka-topics \  
  --bootstrap-server kafka-1:19092 \  
  --create \  
  --topic two-p-topic \  
  --partitions 2 \  
  --replication-factor 1
```

2. You can use the **kafka-topics** tool to describe the details of the topic:

```
$ kafka-topics \  
  --bootstrap-server kafka-1:19092 \  
  --describe \  
  --topic two-p-topic
```

Which will result in this output:

```
Topic: two-p-topic PartitionCount: 2 ReplicationFactor: 1
Configs: segment.bytes=536870912,retention.bytes=536870912
    Topic: two-p-topic Partition: 0 Leader: 101 Replicas: 101
Isr: 101 Offline:
    Topic: two-p-topic Partition: 1 Leader: 103 Replicas: 103
Isr: 103 Offline:
```

- 
- Your Partitions might be placed on different Brokers.
 - The command output when using **--describe** includes configuration settings that are different from Kafka defaults.
 - The custom settings for **segment.bytes** and **retention.bytes** were set on the Lab environment Kafka brokers to limit the amount of message data that is retained and prevent disk full errors from occurring.

3. Use the command-line Producer to write several lines of data to the Topic.

```
$ seq 1 20 > numlist &&
kafka-console-producer \
--bootstrap-server kafka-1:19092 \
--sync \
--topic two-p-topic < numlist
```



The **--sync** option causes the producer to send messages to brokers synchronously, one at a time as they arrive. Using **--sync** in this exercise will result in a better illustration of how the default partitioner assigns records to a partition when they have a **null** key.

4. Use the command-line Consumer to read the messages written to partitions **0** and **1**. Press **Ctrl+C** to exit the Consumer.

```
$ kafka-console-consumer \
  --bootstrap-server \
  kafka-1:19092 \
  --from-beginning \
  --partition 0 \
  --topic two-partition-topic
2
4
6
8
10
12
14
16
18
20
```

```
$ kafka-console-consumer \
  --bootstrap-server \
  kafka-1:19092 \
  --from-beginning \
  --partition 1 \
  --topic two-partition-topic
1
3
5
7
9
11
13
15
17
19
```

5. Note the order of the numbers. What do you think is happening as each message is being produced?
6. Using the previous steps as a guide, create a new Topic with three Partitions, produce the same records into that Topic, and then consume each of the three partitions.

You should see results similar to the following when you consume the three partitions:

```
$ kafka-console-consumer \
  --bootstrap-server \
    kafka-1:19092 \
      --from-beginning \
    \
      --partition 0 \
        --topic three-p-topic
5
7
10
13
16
18
```

```
$ kafka-console-consumer \
  --bootstrap-server \
    kafka-1:19092 \
      --from-beginning \
    \
      --partition 1 \
        --topic three-p-topic
1
3
8
12
15
19
```

```
$ kafka-console-consumer \
  --bootstrap-server \
    kafka-1:19092 \
      --from-beginning \
    \
      --partition 2 \
        --topic three-p-topic
2
4
6
9
11
14
17
20
```

When using the default partitioner and producing records with a `null` key, the first record is randomly assigned to a partition and subsequent records are assigned to the same partition until the batch is flushed to the broker. The next record is then randomly assigned to one of the other partitions and the process repeats itself. When records were produced to the **two-p-topic** with the `--sync option`, the producer randomly selects the next partition to produce each subsequent record to but since there is only one available partition other than the partition just produced to, the records end up being alternately produced to the two available partitions.

With the **three-p-topic** and its three partitions, the producer follows the same random selection of the next partition to produce each subsequent record to, but in this case there are two partitions available and either might get selected. The result is that the records don't get produced to partitions **0**, **1**, and **2** in a repetitive order but instead in a random order with quite possibly an uneven distribution of records across the three partitions.

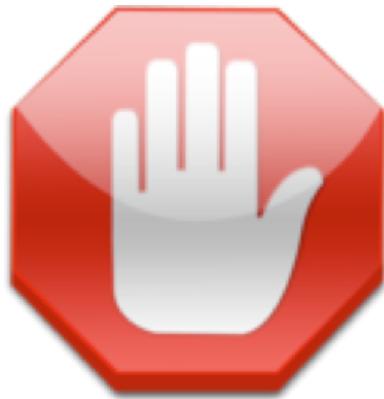
Cleanup

1. Execute the following command to completely cleanup your environment:

```
$ docker-compose down -v
```

Conclusion

You created a Topic with multiple Partitions. You then used the **kafka-console-producer** to write some data to the Topic. Finally you analyzed the order of the data output by the **kafka-console-consumer** and noticed that the order is not deterministic.



STOP HERE. THIS IS THE END OF THE EXERCISE.

Lab 04 Providing Durability

a. Investigating the Distributed Log

In this exercise, you will investigate the distributed log. You will then simulate a Broker failure, and see how to recover the Broker.

Prerequisites

1. Please make sure you have prepared your environment by following → [Preparing the Labs](#)
2. Make sure your Kafka cluster is started, **otherwise** execute this command:

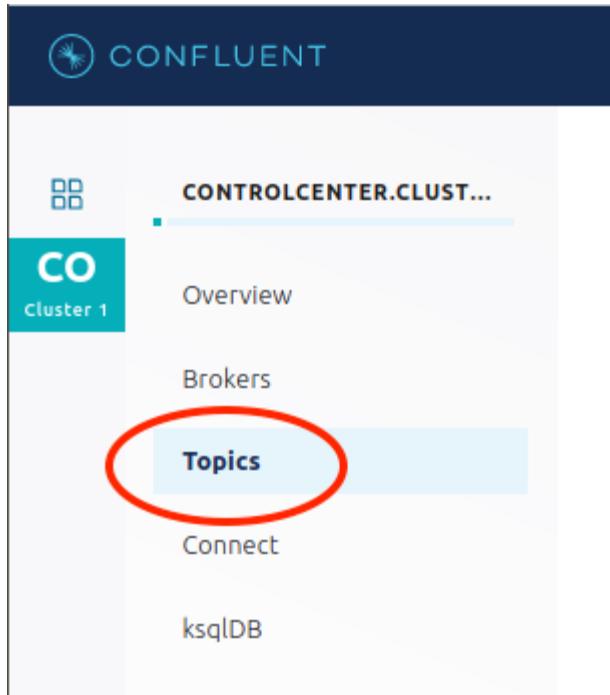
```
$ cd ~/confluent-admin  
$ docker-compose up -d
```

Observing the Distributed Log

1. Create a new Topic called **replicated-topic** with six Partitions and two replicas:

```
$ kafka-topics \  
  --create \  
  --bootstrap-server kafka-1:19092 \  
  --topic replicated-topic \  
  --partitions 6 \  
  --replication-factor 2
```

2. View the Topic information to see that it has been created correctly. The exact output may vary depending on which Brokers the Topic-Partition replicas were assigned.
 - a. Connect to Control Center. If necessary, open a browser tab to the URL <http://localhost:9021>
 - b. In Control Center, click on **Topics**:



- c. Scroll through the Topics list until you find **replicated-topic** and click on the topic name.

Topics	Availability	
Topic name	Under replicated partitions	Out of sync
<u>replicated-topic</u> 	0 of 6	0 of 12
<u>two-p-topic</u>	0 of 2	0 of 2

A detailed view of the topic opens in the right pane with the **Overview** tab selected.

- d. Notice the Topic's Partition list and the corresponding Brokers where each Partition replica resides.

Partitions		Replica placement		Offset	
Partition id	Status	Leader (broker ID)	Followers (broker IDs)	Start	End
0	Available	102	101	0	0
1	Available	101	103	0	0
2	Available	103	102	0	0
3	Available	102	103	0	0
4	Available	101	102	0	0
5	Available	103	101	0	0

- e. You can also view the same Topic information via Kafka command line tools. In your terminal, describe the Topic:

```
$ kafka-topics \
  --describe \
  --bootstrap-server kafka-1:19092 \
  --topic replicated-topic
```

The output should be similar to this:

Topic:replicated-topic	PartitionCount:6	ReplicationFactor:2
Configs:		
Topic: replicated-topic Partition: 0		Leader: 102 Replicas:
102,101 Isr: 102,101		
Topic: replicated-topic Partition: 1		Leader: 103 Replicas:
103,102 Isr: 103,102		
Topic: replicated-topic Partition: 2		Leader: 101 Replicas:
101,103 Isr: 101,103		
Topic: replicated-topic Partition: 3		Leader: 102 Replicas:
102,103 Isr: 102,103		
Topic: replicated-topic Partition: 4		Leader: 103 Replicas:
103,101 Isr: 103,101		
Topic: replicated-topic Partition: 5		Leader: 101 Replicas:
101,102 Isr: 101,102		

3. Produce some data to send to the Topic **replicated-topic**. Leave this process running until it finishes. The following command line sends 6000 messages, 100 bytes in size, at a rate of 1000 messages/sec:

```
$ kafka-producer-perf-test \
  --topic replicated-topic \
  --num-records 6000 \
  --record-size 100 \
  --throughput 1000 \
  --producer-props bootstrap.servers=kafka-1:19092,kafka-
2:29092,kafka-3:39092
```

The output should look similar to this:

```
4999 records sent, 999.8 records/sec (0.10 MB/sec), 5.2 ms avg
latency, 298.0 max latency.
6000 records sent, 998.336106 records/sec (0.10 MB/sec), 4.69 ms avg
latency, 298.00 ms max latency, 2 ms 50th, 25 ms 95th, 35 ms 99th, 46
ms 99.9th.
```

4. Start the console Producer for the same Topic **replicated-topic**:

```
$ kafka-console-producer \
  --bootstrap-server kafka-1:19092,kafka-2:29092,kafka-3:39092 \
  --topic replicated-topic
```

5. At the > prompt, type "I heart logs" and press Enter. Add five more messages and then press **Ctrl+D** to exit the console Producer:

```
> I heart logs
> Hello world
> All your base
> Kafka rules
> Don't worry
> Be happy
<Ctrl+D>
```

Examine Checkpoint Files

For each Broker (**kafka-1**, **kafka-2** and **kafka3**) examine the contents of the checkpoint files **recovery-point-offset-checkpoint** and **replication-offset-checkpoint** at **/var/lib/kafka/data**. Let's start with the first Broker called **kafka-1**:

1. The next steps will be run on the Broker. From your host system, open a new terminal and connect to the Broker, here **kafka-1**:

```
$ cd ~/confluent-admin  
$ docker-compose exec kafka-1 /bin/bash  
[appuser@kafka-1 ~]$
```

2. The **recovery-point-offset-checkpoint** file records the point up to which data has been flushed to disks. This is important as, on hard failure, the Broker needs to scan unflushed data, verify the CRC, and truncate the corrupted log.

```
[appuser@kafka-1 ~]$ cat \  
/var/lib/kafka/data/recovery-point-offset-checkpoint | \  
grep replicated-topic  
replicated-topic 1 0  
replicated-topic 2 0  
replicated-topic 3 0  
replicated-topic 4 0
```

In the output above, the first number is the Partition number, and the second number is the offset of the last flushing point. In the output above, the offset of the last flushing point is expected to be 0 because the open file segment has not been flushed yet

3. The **replication-offset-checkpoint** file is the offset of the last committed offset (e.g. the **high water mark**).

```
[appuser@kafka-1 ~]$ cat \  
/var/lib/kafka/data/replication-offset-checkpoint | \  
grep replicated-topic  
replicated-topic 3 922  
replicated-topic 1 1077  
replicated-topic 2 1078  
replicated-topic 4 1133
```

In the output above, the first number is the Partition number, and the second number is the offset of the high water mark. The offset of the high water mark is expected to be different on each partition because the number of produced messages, 6006 (6000 + 6) was not evenly distributed across the six Partitions due to the sticky partitioner.



Your high water marks will likely be slightly different than those shown in the above example.

4. Disconnect from the Broker container by pressing **Ctrl+D**.

5. Repeat the previous steps for the other two Brokers **kafka-2** and **kafka-3**.

Examining Topic Partitions

On each of the Brokers, examine the contents of one of the data directories for at least one of the Topic-Partitions. For example, the directory **replicated-topic-0** has log files for Partition 0 of Topic **replicated-topic**. The following steps will list the log files for Partition 0 on the **kafka-1** Broker.

1. From your host system connect to the **kafka-1** Broker container:

```
$ cd ~/confluent-admin
$ docker-compose exec kafka-1 /bin/bash
[appuser@kafka-1 ~]$
```

2. List the directory **replicated-topic-0**:

```
[appuser@kafka-1 ~]$ ls -l /var/lib/kafka/data/replicated-topic-0
total 132
-rw-r--r-- 1 appuser appuser 10485760 Oct 23 01:35
000000000000000000000000.index
-rw-r--r-- 1 appuser appuser    121393 Oct 23 01:36
000000000000000000000000.log
-rw-r--r-- 1 appuser appuser 10485756 Oct 23 01:35
000000000000000000000000.timeindex
-rw-r--r-- 1 appuser appuser          8 Oct 23 01:34 leader-epoch-
checkpoint
```

The command may return the following error:

```
ls: cannot access /var/lib/kafka/data/replicated-topic-0:  
No such file or directory
```

Why would this error occur?

The directory **replicated-topic-0** does not exist on all three Brokers in your cluster.



Why not?

Remember that a Broker may not store every Partition for a Topic. As shown in the Partition lists earlier, each Broker only contains a subset of the Partitions for **replicated-topic**. If you get an error when trying to list the **replicated-topic-0** directory, the Partition is on another Broker.

If this error occurs, repeat the command using **replicated-topic-1**.

The data files have the following meaning:

- **.log** holds the messages and metadata
- **.index** maps message offsets to their byte position in the log file
- **.timeindex** maps message timestamps to their offset number
- **leader-epoch-checkpoint** maps the leader epoch to its corresponding start offset

3. On the Brokers, view the contents of the file **leader-epoch-checkpoint**:

```
[appuser@kafka-1 ~]$ cat \  
/var/lib/kafka/data/replicated-topic-0/leader-epoch-checkpoint
```

This file tracks the lineage of the Partition leadership by noting which offset starts a new leader generation. The first row denotes the version of the file format, in this case **0**. The second row denotes the number of epochs, which is incremented every leader election. This Partition has only had one leader in its history, so it is still in its first epoch. Each row afterwards shows the epoch ID followed by the offset where the epoch begins. In this

case, the current epoch ID is **0** and starts at offset **0**.

```
0  
1  
0 0
```

4. On the Broker, use the **kafka-dump-log** tool to view details of the **.log** file:

```
[appuser@kafka-1 ~]$ kafka-dump-log \  
--print-data-log \  
--files \  
/var/lib/kafka/data/replicated-topic-0/000000000000000000000000.log
```

The command output should be similar to (shortened):

```
Dumping /var/lib/kafka/data/replicated-topic-  
0/000000000000000000000000.log  
Starting offset: 0  
baseOffset: 0 lastOffset: 1 count: 2 baseSequence: -1 lastSequence:  
-1 producerId: -1 producerEpoch: -1 partitionLeaderEpoch: 0  
isTransactional: false isControl: false position: 0 CreateTime:  
1603416922512 size: 279 magic: 2 compresscodec: NONE crc: 2708005014  
isValid: true  
| offset: 0 CreateTime: 1603416922512 keysize: -1 valuesize: 100  
sequence: -1 headerKeys: [] payload:  
SSXVNJHPDQDXVCRASTVYBCWVMGNYKRXVZXKGXTSPSJGYLUEGQFLAQLOCFJBEPOWFNSO  
MYARHAOPUFOJHHDXEHXJBHWGSMZJGNL  
| offset: 1 CreateTime: 1603416922512 keysize: -1 valuesize: 100  
sequence: -1 headerKeys: [] payload:  
SSXVNJHPDQDXVCRASTVYBCWVMGNYKRXVZXKGXTSPSJGYLUEGQFLAQLOCFJBEPOWFNSO  
MYARHAOPUFOJHHDXEHXJBHWGSMZJGNL  
...  
| offset: 887 CreateTime: 1603416928132 keysize: -1 valuesize: 100  
sequence: -1 headerKeys: [] payload:  
SSXVNJHPDQDXVCRASTVYBCWVMGNYKRXVZXKGXTSPSJGYLUEGQFLAQLOCFJBEPOWFNSO  
MYARHAOPUFOJHHDXEHXJBHWGSMZJGNL  
baseOffset: 888 lastOffset: 888 count: 1 baseSequence: -1  
lastSequence: -1 producerId: -1 producerEpoch: -1  
partitionLeaderEpoch: 0 isTransactional: false isControl: false  
position: 121314 CreateTime: 1603416972594 size: 79 magic: 2  
compresscodec: NONE crc: 4111803574 isValid: true  
| offset: 888 CreateTime: 1603416972594 keysize: -1 valuesize: 11  
sequence: -1 headerKeys: [] payload: Don't worry
```

5. Questions:

- a. Are the offsets in the **.log** file sequential?

- b. Can you work out what is the offset of the message "All your base" that you entered earlier? If you can't find the value in the directory **replicated-topic-0**, where else might it be?
6. On the Broker, use the **DumpLogSegments** tool to view the offsets in the **.index** file:

```
[appuser@kafka-1 ~]$ kafka-run-class kafka.tools.DumpLogSegments \
--files \
/var/lib/kafka/data/replicated-topic-0/00000000000000000000000000000000.index

Dumping /var/lib/kafka/data/replicated-topic-
0/00000000000000000000000000000000.index
Dumping /var/lib/kafka/data/replicated-topic-
0/00000000000000000000000000000000.index
offset: 32 position: 4172
offset: 64 position: 8466
offset: 94 position: 12638
...
...
offset: 807 position: 109936
offset: 837 position: 114060
offset: 867 position: 118245
```



Note the difference from one index entry position to the next is ~4000 which corresponds to the default **log.index.interval.bytes=4096** setting.

7. Disconnect from the Broker (container) by pressing **Ctrl+D**.
8. Repeat the above steps for the other two Brokers **kafka-2** and **kafka-3**.

Taking a Broker Offline

1. First let's define an **action** in Control Center for when an under replication happens. We will see later why that is important:
 - a. Open Control Center
 - b. Click the bell icon in the top right corner to navigate to **Alerts**
 - c. Click the button **Create trigger**
 - d. Complete the dialog using the following settings:

Field	Value
Trigger Name	Under Replicated Partitions
Component type	Cluster
Cluster id	controlcenter.cluster ...
Metric	Under replicated topic partitions
Condition	Greater than
Value	0

New trigger

General

Trigger name* _____

Under Replicated Partitions

Components

Component type* _____

Cluster

Cluster id* _____

controlcenter.cluster [bnZHH_GSRdyOyN0zscsS2Q] ×

Criteria

Metric* _____

Under replicated topic partitions

Condition* _____

Greater than

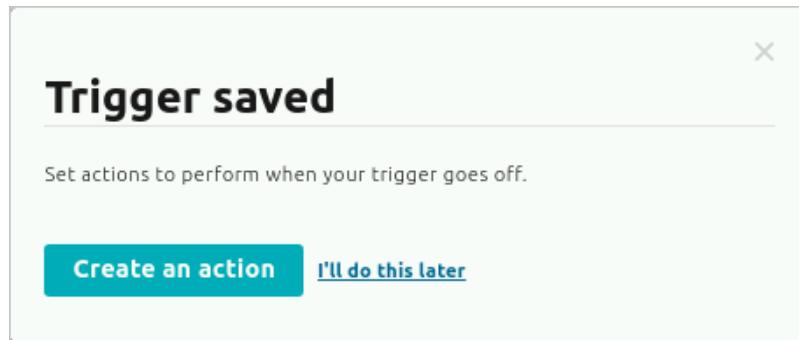
Value* _____

0

Save**Cancel**

e. Click **Save**

f. In the confirmation box, when asked to "Set actions to perform when your trigger goes off", select **Create an action**



g. Complete the dialog using the following settings:

Field	Value
Action Name	My action
Triggers	Under Replicated Partitions
Actions	Send email
Subject	Under replicated partitions
Max send rate	1 Per minute
Recipient email address	anything@anything.com

New action

General

Action Name* —

My action

Enabled

Triggers

Triggers* —

Under Replicated Partitions x

Actions

Action* —

Send email

Subject* —

Under replicated partitions

Max send rate* —

1

Per minute

Recipient email address* —

anything@anything.com

Save

Cancel

h. Click **Save**



This lab will not actually trigger emails since email/SMTP settings have not been configured.

- From a terminal window discover which of the 3 Brokers currently is the controller by using this command:

```
$ zookeeper-shell zk-1:12181 get /controller
Connecting to zk-1:12181

WATCHER:::

WatchedEvent state:SyncConnected type:None path:null
{"version":1,"brokerid":101,"timestamp":"1603416838481"}
```

In this particular case, Broker **kafka-1** with brokerid **101** is the Controller.

- To capture the Controller Broker name in an environment variable, execute the command that corresponds to your Controller **brokerid**:

Controller	Command
101	\$ export INITIAL_CONTROLLER=kafka-1
102	\$ export INITIAL_CONTROLLER=kafka-2
103	\$ export INITIAL_CONTROLLER=kafka-3

- Let's simulate the worst case. Stop the Broker that is acting as the Controller to see what happens. From your host system execute:

```
$ cd ~/confluent-admin
$ docker-compose stop $INITIAL_CONTROLLER

Stopping kafka-1 ... done
```

- Verify that the other Brokers are running:

- From your host system use the following command:

\$ docker-compose ps	Name	Command	State
Ports			
\-----			
connect	/etc/confluent/docker/run	Up (healthy)	
0.0.0.0:8083->8083/tcp, 9092/tcp			
control-center	/etc/confluent/docker/run	Up	
0.0.0.0:9021->9021/tcp			
kafka-1	/etc/confluent/docker/run	Exit 143	
kafka-2	/etc/confluent/docker/run	Up	
0.0.0.0:29092->29092/tcp, 9092/tcp			
kafka-3	/etc/confluent/docker/run	Up	
0.0.0.0:39092->39092/tcp, 9092/tcp			
ksqldb-server	/etc/confluent/docker/run	Up (healthy)	
0.0.0.0:8088->8088/tcp			
schema-registry	/etc/confluent/docker/run	Up	
0.0.0.0:8081->8081/tcp			
tools	/bin/bash	Up	
zk-1	/etc/confluent/docker/run	Up	
0.0.0.0:12181->12181/tcp, 2181/tcp, 2888/tcp, 3888/tcp			
zk-2	/etc/confluent/docker/run	Up	
2181/tcp, 0.0.0.0:22181->22181/tcp, 2888/tcp, 3888/tcp			
zk-3	/etc/confluent/docker/run	Up	
2181/tcp, 2888/tcp, 0.0.0.0:32181->32181/tcp, 3888/tcp			

- b. Make sure that the **State** of the other Brokers marked as **Up** and initial Controller Broker as **Exit 143**.
- c. From a terminal window use **zookeeper-shell** command-line tool to see which Brokers are registered with ZooKeeper:

```
$ zookeeper-shell zk-1:12181 ls /brokers/ids

Connecting to zk-1:2181

WATCHER:::

WatchedEvent state:SyncConnected type:None path:null
[102, 103]
```

In our example, the command output shows that only Brokers **102** and **103** are active.

- 6. From the host system, review the server log for the initial Controller. Look for confirmation of the controlled shutdown succeeding:

```
$ docker-compose logs $INITIAL_CONTROLLER | tail  
...  
kafka-1 | [2020-10-23 01:52:12,916] INFO [Producer  
clientId=confluent-metrics-reporter] Proceeding to force close the  
producer since pending requests could not be completed within timeout  
0 ms. (org.apache.kafka.clients.producer.KafkaProducer)  
kafka-1 | [2020-10-23 01:52:12,920] INFO [KafkaServer  
id=101] shut down completed (kafka.server.KafkaServer)
```



All logs produced by the Brokers running inside a container are written to STDOUT and STDERR and reflected in the Docker logs. On a dedicated Broker machine, this server log data will be located at [/var/log/kafka/server.log](#).

7. From the host system, review the Controller log for the initial Controller. Look for confirmation of the Controller role being released:

```
$ docker-compose logs $INITIAL_CONTROLLER | grep Resigned  
kafka-1 | [2020-10-23 01:52:10,707] INFO [Controller  
id=101] Resigned (kafka.controller.KafkaController)
```

8. **Question:** If you don't find the above entry in the log of the initial Controller Broker, where else could it be? Could another Broker have been the active controller?
9. Wait up to five minutes and then observe the cluster in Control Center.
 - In Control Center click **Cluster 1**.
 - In the **Broker** panel of **Overview**, observe the **Total** count decrease to **2**:
 - In the **Topics** panel of **Overview**, observe the **Under replicated partitions** count (should be non-zero).
 - Click on **Topics**.
 - Scroll through the topic list to locate **replicated-topic** and click the topic name.
 - On the **Overview** tab, notice there are now **Under replicated partitions**:

[ALL TOPICS >](#)

replicated-topic

[Overview](#) [Messages](#) [Schema](#) [Configuration](#)

Production

0

Bytes per second



Consumption

--

Bytes per second

Availability

4 of 6

Under replicated partitions



4 of 12

Out of sync followers



- g. Click the bell icon in the top right corner to view the alert history and notice that the broker down event triggered an alert. It is critical to monitor under replicated partitions to ensure message durability in your Kafka cluster:

[ALERTS >](#)

Overview

[History](#) [Triggers](#) [Actions](#) [REST API](#)

Time	Trigger	Component	Action Count
a few seconds ago	Under Replicated Partitions	pMs6rSQEScqBw5AudfLvug	1

- h. Start the console Producer for the same Topic **replicated-topic**:

```
$ kafka-console-producer \
  --bootstrap-server kafka-1:19092,kafka-2:29092,kafka-3:39092 \
  --topic replicated-topic
```



Clients send the cluster metadata request to **bootstrap-server** brokers in a random order. You may see a warning message if the initial metadata request is sent to a broker that is not running.

- i. At the > prompt, type six more messages and then press **Ctrl+D** to exit the console Producer:

```
> Kafka
> Distributed
> Secure
> Real-time
> Scalable
> Fast
<Ctrl+D>
```

10. From your host system view the impact to leader epoch:

- a. First run **kafka-topics** again to identify which **replicated-topic** partitions had a preferred replica that was the initial Controller Broker. It would have been the leader replica when the topic was created. Since we shut this broker down, leader election would have occurred and these partitions would have been assigned a new leader:

```
$ kafka-topics \
  --describe \
  --bootstrap-server kafka-1:19092,kafka-2:29092,kafka-3:39092 \
  --topic replicated-topic
```

The output should be similar to this:

```
Topic: replicated-topic PartitionCount: 6 ReplicationFactor: 2
Configs: segment.bytes=536870912,retention.bytes=536870912
    Topic: replicated-topic Partition: 0 Leader: 102 Replicas:
102,103 Isr: 102,103 Offline:
    Topic: replicated-topic Partition: 1 Leader: 102 Replicas:
101,102 Isr: 102 Offline: 101
    Topic: replicated-topic Partition: 2 Leader: 103 Replicas:
103,101 Isr: 103 Offline: 101
    Topic: replicated-topic Partition: 3 Leader: 102 Replicas:
102,101 Isr: 102 Offline: 101
    Topic: replicated-topic Partition: 4 Leader: 103 Replicas:
101,103 Isr: 103 Offline: 101
    Topic: replicated-topic Partition: 5 Leader: 103 Replicas:
103,102 Isr: 103,102 Offline:
```

- b. On either of the other two Brokers, view the contents of the file **leader-epoch-checkpoint** in one of the Topic-Partitions subdirectory for which the initial Controller Broker was the leader replica:

```
$ docker-compose exec kafka-2 cat \
/var/lib/kafka/data/replicated-topic-2/leader-epoch-checkpoint
0
2
0 0
1 1077
```

This Partition now reflects that it has had two leaders. Leader epoch **0** had an initial offset of **0**. Leader epoch **1** has an initial offset of **1077**.

- 
 - Your leader epoch **1** may have a slightly different initial offset.
 - If you see the following response to the above command, it indicates that no leader election has occurred for that partition.

```
0
1
0 0
```

Bringing a Broker Online

Bringing a Broker online is as simple as restarting the Broker and letting Kafka automatically rebalance the leaders.

1. From your host system restart the initial Controller Broker:

```
$ docker-compose start $INITIAL_CONTROLLER
```

2. Wait five minutes and then observe the cluster in Control Center:

- In Control Center, click **Cluster 1**.
- In the **Broker** panel of **Overview**, observe the **Total** count as it returns to **3**.
- In the **Topics** panel of **Overview**, observe the **Under replicated partitions** count as it returns to **0**.
- Click on **Topics**.
- Scroll through the topic list to locate **replicated-topic** and click the topic name.
- On the **Overview** tab, notice the **Under replicated partitions** has returned to **0**.



STOP HERE. THIS IS THE END OF THE EXERCISE.

Lab 05 Configuring a Kafka Cluster

a. Exploring Configuration

In this exercise, you will research a few important Broker properties of a Kafka cluster.

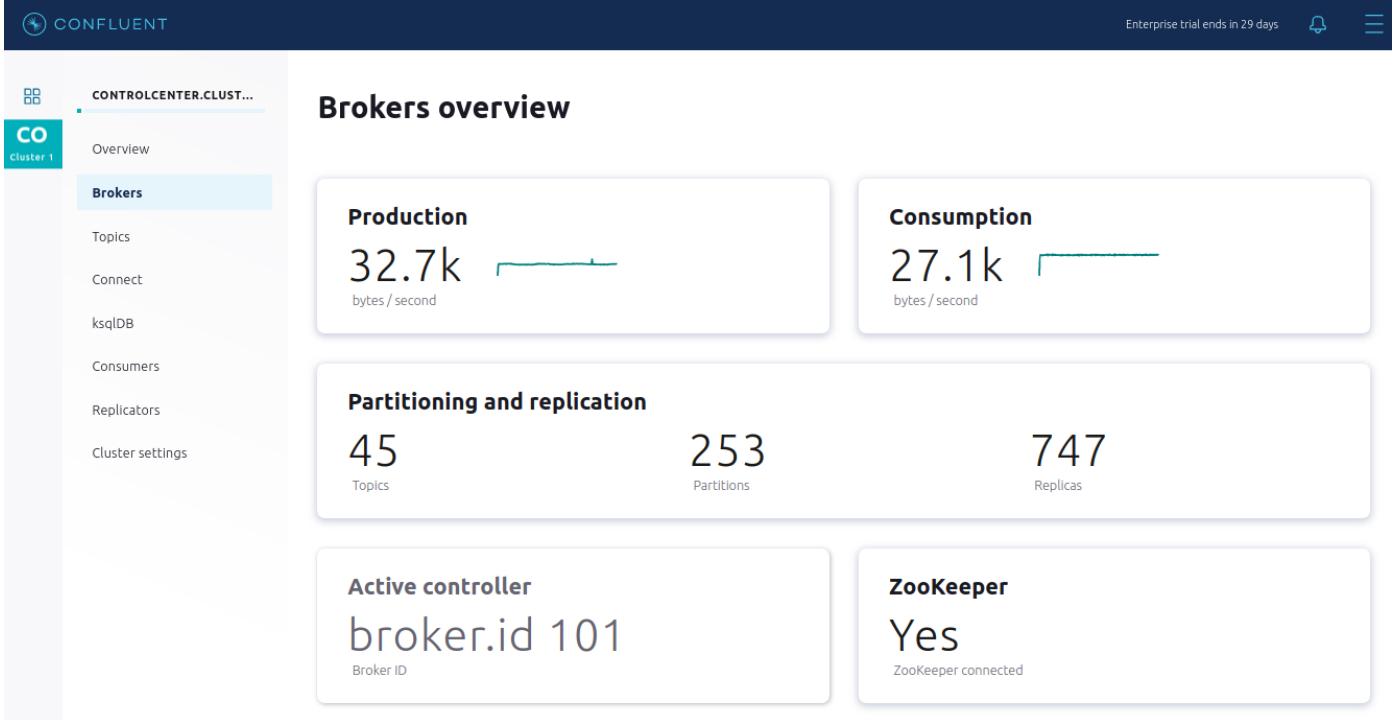
Prerequisites

1. Please make sure you have prepared your environment by following → [Preparing the Labs](#)
2. Make sure your Kafka cluster is started, **otherwise** execute this command:

```
$ cd ~/confluent-admin  
$ docker-compose up -d
```

Configuration Jigsaw

1. Connect to Control Center. If necessary, open a browser tab to the URL <http://localhost:9021>.
2. Select **Cluster 1** on the left hand side, then click **Brokers** and you will see the **Brokers overview** view:



Observe the metrics you see on this view. Do they correspond to what you'd expect to view on a first glance? Why yes? Why not?

3. On the left side select **Cluster settings** and then select the tab **Broker defaults**:

Cluster settings

Broker defaults ←

Hide settings set to default value Edit Settings Download

GENERAL

- broker.id 3 different values*

• broker.101	101
• broker.102	102
• broker.103	103

LISTENER

- listeners 3 different values*

• broker.101	DOCKER://kafka-1:9092,HOST://kafka-1:1...
• broker.102	DOCKER://kafka-2:9092,HOST://kafka-2:2...
• broker.103	DOCKER://kafka-3:9092,HOST://kafka-3:3...

- inter.broker.listener.name DOCKER
- advertised.listeners 3 different values*

• broker.101	DOCKER://kafka-1:9092,HOST://kafka-1:1...
• broker.102	DOCKER://kafka-2:9092,HOST://kafka-2:2...
• broker.103	DOCKER://kafka-3:9092,HOST://kafka-3:3...

LOG

- log.dirs /var/lib/kafka/data

- Click the **Hide settings set to default value** toggle button to display all Cluster settings in the list.
- Spend 5 minutes looking at the different categories of properties, and perhaps read a few of the descriptions provided (the description pops up when hovering your mouse over a property). There are many configuration settings, so don't get too focused on one property in particular.
 - What do you notice? Write down a few observations.
 - What do you wonder? Write down a few questions that arise.
- Here are some sets of configuration properties. Choose **one** set of properties and take 15

minutes to research what they do, what some reasonable values are, and what some unreasonable (but valid) values might be.

- listeners:

- `listeners`
- `advertised.listeners`
- `listener.security.protocol.map`
- `control.plane.listener.name`
- `inter.broker.listener.name`

- logs (i.e. cluster default properties for Topics, Partitions, and replicas):

- `log.dirs`
- `default.replication.factor`
- `log.retention.hours`
- `log.segment.bytes`
- `log.roll.ms`
- `num.partitions`
- `replica.lag.time.max.ms`

- network:

- `broker.id`
- `zookeeper.connect`
- `zookeeper.session.timeout.ms`
- `group.max.session.timeout.ms`
- `sasl.enabled.mechanisms`
- `max.connections.per.ip`
- `max.connections.per.ip.overrides`

- CPU:

- `num.io.threads`

- `num.network.threads`
- `num.recovery.threads.per.data.dir`
- `num.replica.fetchers`
- `background.threads`
- `log.cleaner.threads`



This is just a first exposure. Many of these properties will be explored in more detail in upcoming material, so don't worry too much if you still have questions at this point. Record your questions and come back to them later towards the end of training to make sure they get answered.



STOP HERE. THIS IS THE END OF THE EXERCISE.

b. Increasing Replication Factor

We all make mistakes. Maybe you accidentally created a mission-critical Topic with a replication factor of 1. Luckily, you notice this before disaster strikes. In this exercise, you will use the **kafka-reassign-partitions** tool to increase the replication factor of a Topic. The **kafka-reassign-partitions** tool is an included utility that is usually used to rebalance the load of Partitions across Brokers. It will be discussed in further detail later in the course.

Prerequisites

1. Please make sure you have prepared your environment by following → [Preparing the Labs](#)
2. Make sure your Kafka cluster is started, **otherwise** execute this command:

```
$ cd ~/confluent-admin  
$ docker-compose up -d
```

Increasing Replication Factor

1. Create a Topic called **test** with 3 Partitions and replication factor 1.

```
$ kafka-topics \  
  --bootstrap-server kafka-1:19092,kafka-2:29092,kafka-3:39092 \  
  --create \  
  --topic test \  
  --partitions 3 \  
  --replication-factor 1
```

2. View the Topic information. Partitions may land on different Brokers than shown here.

```
$ kafka-topics \
  --bootstrap-server kafka-1:19092,kafka-2:29092,kafka-3:39092 \
  --describe \
  --topic test

Topic:test PartitionCount:3      ReplicationFactor:1 Configs:
  Topic: test Partition: 0        Leader: 103 Replicas: 103    Isr: 103
  Topic: test Partition: 1        Leader: 102 Replicas: 102    Isr: 102
  Topic: test Partition: 2        Leader: 101 Replicas: 101    Isr: 101
```

3. Create a **json** file called **replicate_topic_test_plan.json** that declares your desired state of Partition replication. Notice that **kafka-reassign-partitions** allows you to set replication factor on a per-Partition basis.

```
$ cat << EOF > replicate_topic_test_plan.json
{"version":1,
 "partitions":[
  {"topic":"test","partition":0,"replicas":[101,102,103]},
  {"topic":"test","partition":1,"replicas":[101,102,103]},
  {"topic":"test","partition":2,"replicas":[102,103]}
]
EOF
```



Notice that we declare Partition 2 to only have 2 replicas. This shows the granular control of the **kafka-reassign-partitions** tool. Also note the first of each list will become the preferred replica.

4. Use the **--reassignment-json-file** and **--execute** options of the **kafka-reassign-partitions** tool to execute the change.

```
$ kafka-reassign-partitions \
  --bootstrap-server kafka-1:19092,kafka-2:29092,kafka-3:39092 \
  --reassignment-json-file replicate_topic_test_plan.json \
  --zookeeper zk-1:12181 \
  --execute
```

Current partition replica assignment

```
{"version":1,"partitions":[{"topic":"test","partition":2,"replicas":[101],"log_dirs":["any"]},{ {"topic":"test","partition":1,"replicas":[102],"log_dirs":["any"]},{ {"topic":"test","partition":0,"replicas":[103],"log_dirs":["any"]}]}}
```

Save this to use as the --reassignment-json-file option during rollback

Successfully started partition reassignments for test-0,test-1,test-2

5. View the Topic information again to see that the Partitions are now replicated.

```
$ kafka-topics \
  --bootstrap-server kafka-1:19092,kafka-2:29092,kafka-3:39092 \
  --describe \
  --topic test
```

```
Topic:test PartitionCount:3      ReplicationFactor:3 Configs:
      Topic: test Partition: 0      Leader: 103 Replicas: 101,102,103
Isr: 103,101,102
      Topic: test Partition: 1      Leader: 102 Replicas: 101,102,103
Isr: 102,101,103
      Topic: test Partition: 2      Leader: 102 Replicas: 102,103   Isr:
102,103
```



STOP HERE. THIS IS THE END OF THE EXERCISE.

Lab 06 Managing a Kafka Cluster

a. Kafka Administrative Tools

In this exercise, you will delete a Topic, reassign Partitions, and simulate a completely failed Broker.

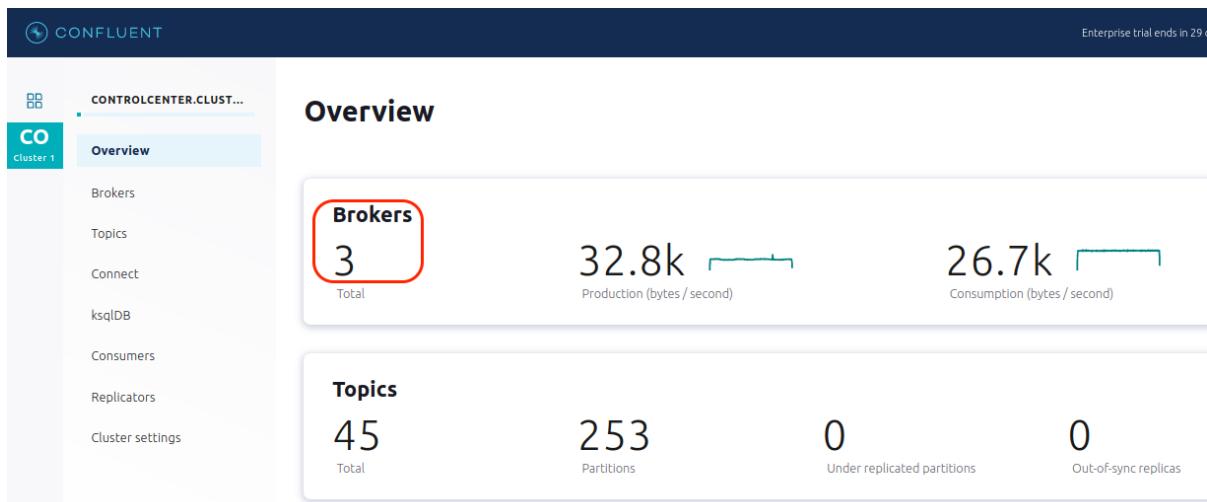
Prerequisites

1. Please make sure you have prepared your environment by following → [Preparing the Labs](#)
2. Make sure your Kafka cluster is started, **otherwise** execute this command:

```
$ cd ~/confluent-admin  
$ docker-compose up -d
```

Deleting Topics in the Cluster

1. Connect to Control Center. If necessary, open a browser tab to the URL <http://localhost:9021>.
2. Verify that three Brokers are running. In the Control Center **Overview** view, observe the Broker count is three.



3. Delete the Topic **replicated-topic**:

```
$ kafka-topics \
  --bootstrap-server kafka-1:19092 \
  --delete \
  --topic replicated-topic
```



When deleting a Topic, always make sure you don't have any Producers or Consumers running. Otherwise, the Topic will be re-created automatically. Also, you may have already deleted this Topic when destroying the cluster with **docker-compose down -v**. If so, create the Topic again, verify it exists, and then delete it.

4. In the Control Center, check that the Topic **replicated-topic** is now gone from the list of topics in the Topic Management view.



The Topic may not disappear immediately. If necessary wait a few minutes.

Rebalancing the Cluster

1. Create a new Topic called **moving** with 6 Partitions and 2 replicas, on only Broker 1 and Broker 2 (with IDs **101** and **102**):

```
$ kafka-topics \
  --bootstrap-server kafka-1:19092 \
  --create \
  --topic moving \
  --replica-assignment
101:102,102:101,101:102,102:101,101:102,102:101
```

2. Run the command line tool **kafka-producer-perf-test** to produce 2GB of data to Topic **moving**.



Wait until this command has completed before continuing.

```
$ kafka-producer-perf-test \
--topic moving \
--num-records 2000000 \
--record-size 1000 \
--throughput 1000000000 \
--producer-props bootstrap.servers=kafka-1:19092,kafka-2:29092
```

3. Verify which Brokers contain Partitions for the Topic **moving**:

- In the Control Center UI make sure you select **Cluster 1**
- Then select the tab **Topics**:

The screenshot shows the Confluent Control Center interface. The top navigation bar has the Confluent logo and the text 'CONFLUENT'. Below it, there's a sidebar with icons for Overview, Brokers, Topics (which is highlighted in blue), Connect, ksqlDB, Consumers, Replicators, and Cluster settings. The main content area is titled 'All topics' and contains a search bar with 'Search topics' and a 'Hide internal topics' checkbox. A table lists topics with columns: 'Topics' (Topic name), 'Availability' (Under replicated partitions, Out of sync followers), and 'Throughput' (Bytes/sec produced). Two topics are listed: 'moving' and 'two-p-topic'. A red arrow points to the 'moving' topic in the table.

Topics	Availability	Throughput	
Topic name	Under replicated partitions	Out of sync followers	Bytes/sec produced
moving	0 of 6	0 of 12	--
two-p-topic	0 of 2	0 of 2	0

- In the list of topics select the **moving** topic:

Production
0 Bytes per second

Consumption
-- Bytes per second

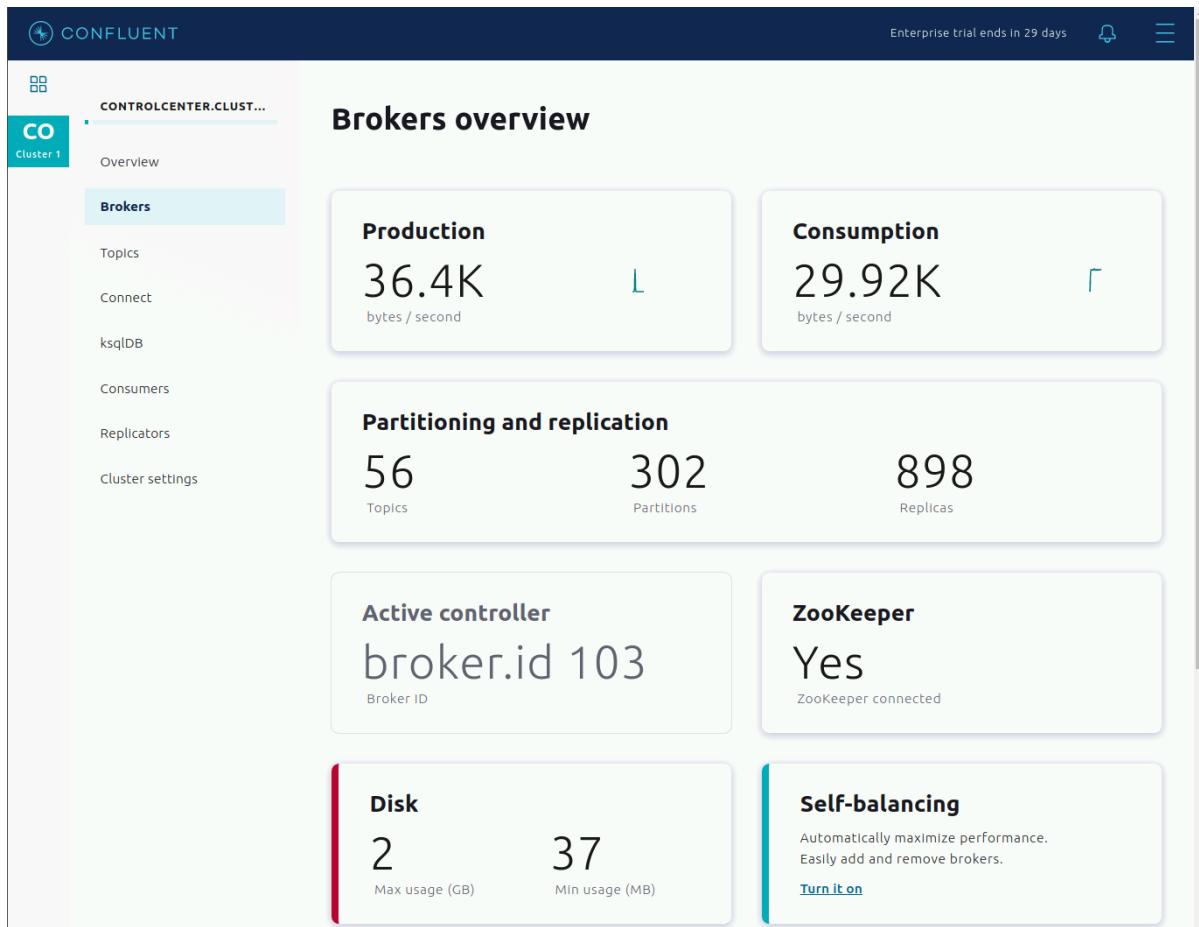
Availability
0 of 6 Under replicated partitions 0 of 12 Out of sync followers

Partitions		Replica placement		Followers (broker IDs)	Offset		Size Total Size
Partition ID	Status	Leader (broker ID)	Follower (broker IDs)		Start	End	
0	Available	101	102		0	333333	338MB
1	Available	102	101		0	333334	338MB
2	Available	101	102		0	333334	338MB
3	Available	102	101		0	333333	338MB
4	Available	101	102		0	333333	338MB
5	Available	102	101		0	333333	338MB



Notice the Topic's Partition list and the corresponding Brokers where each Partition replica resides. They are all on Brokers 101 and 102.

4. In Control Center on the left side, select **Brokers**.
5. In the **Brokers overview** view, look at the metric **Disk**. Notice the red bar to indicate a cluster imbalance: (may take a couple of minutes)



6. Consume some data from the Topic **moving**. Leave this consumer running for the duration of the exercise:

```
$ kafka-consumer-perf-test \
--bootstrap-server kafka-1:19092,kafka-2:29092 \
--topic moving \
--group test-group \
--threads 1 \
--show-detailed-stats \
--timeout 1000000 \
--reporting-interval 5000 \
--messages 10000000
```

```
time, threadId, data.consumed.in.MB, MB.sec, data.consumed.in.nMsg, nMsg.sec, rebalance.time.ms, fetch.time.ms, fetch.MB.sec, fetch.nMsg.sec  
2018-10-03 12:21:58:622, 0, 172.8697, 34.1437, 181267, 35802.2911, 3033, 2030, 85.1575, 89294.0887  
2018-10-03 12:22:03:622, 0, 684.5884, 102.3438, 717843, 107315.2000, 0, 5000, 102.3438, 107315.2000  
2018-10-03 12:22:08:622, 0, 1249.4726, 112.9768, 1310167, 118464.8000, 0, 5000, 112.9768, 118464.8000  
2018-10-03 12:22:13:624, 0, 1646.5025, 79.3742, 1726483, 83229.9080, 0, 5002, 79.3742, 83229.9080  
...
```

The next few steps use the Confluent Auto Data Balancer to rebalance the cluster. If you prefer to use the Apache open source tool instead, you may view the Appendix section of this manual called **Reassigning Partitions in a Topic: Alternate Method** (→ [Jump here](#)).



Note the open source tool requires users to enumerate every Topic to rebalance and does not take into consideration disk utilization per Broker.

7. Run the Confluent Auto Data Balancer to rebalance the cluster. This will distribute the Partitions in the cluster, including the Topic **moving**, so that the Brokers are more evenly utilized. Wait for a few seconds for it to compute the rebalance plan.

```
$ confluent-rebalancer execute \  
  --bootstrap-server kafka-1:19092 \  
  --metrics-bootstrap-server kafka-1:19092,kafka-2:29092 \  
  --throttle 1000000 \  
  --verbose
```

Computing the rebalance plan (this may take a while) ...
 You are about to move 5 replica(s) for 5 partitions to 2 broker(s) with total size 1,348.5 MB.
 The preferred leader for 9 partition(s) will be changed.
 In total, the assignment for 10 partitions will be changed.
 The minimum free volume space is set to 20.0%.

The following brokers will have less than 40% of free volume space during the rebalance:

Broker During Rebalance	Current Size (MB)	Size During Rebalance (MB)	Free %
	Size After Rebalance (MB)	Free % After	
101	2,046.8	2,046.8	36.2
2,046.8		36.2	
103	21.2	1,369.7	31.8
1,369.7		31.8	
102	2,046.8	2,046.8	36.2
2,046.8		36.2	

Min/max stats for brokers (before -> after):

Type	Leader Count	Replica Count	Size (MB)
Min	99 (id: 101) -> 101 (id: 101)	298 (id: 103) -> 301 (id: 101)	
21.2 (id: 103) -> 1,369.7 (id: 103)			
Max	105 (id: 103) -> 102 (id: 102)	303 (id: 101) -> 301 (id: 101)	
2,046.8 (id: 101) -> 1,374 (id: 102)			

No racks are defined.

Broker stats (before -> after):

Broker Space (%)	Leader Count	Replica Count	Size (MB)	Free
101 -> 38.3	99 -> 101	303 -> 301	2,046.8 -> 1,371	36.2
102 -> 38.3	101 -> 102	302 -> 301	2,046.8 -> 1,374	36.2
103 -> 31.8	105 -> 102	298 -> 301	21.2 -> 1,369.7	36.2

Would you like to continue? (y/n):

8. To start the rebalance operation, type **y** and press **Enter**.

The following response will appear:

The rebalance has been started, run `status` to check progress.

Warning: You must run the `status` or `finish` command periodically, until the rebalance completes, to ensure the throttle is removed. You can also alter the throttle by re-running the execute command passing a new value.

9. Observe the configured throttling limits:

```
$ kafka-configs \
  --bootstrap-server kafka-1:19092 \
  --describe \
  --entity-type brokers
Dynamic configs for broker 101 are:
  leader.replication.throttled.rate=1000000 sensitive=false
  synonyms={DYNAMIC_BROKER_CONFIG:leader.replication.throttled.rate=100000}
  follower.replication.throttled.rate=1000000 sensitive=false
  synonyms={DYNAMIC_BROKER_CONFIG:follower.replication.throttled.rate=1000000}
Dynamic configs for broker 102 are:
  leader.replication.throttled.rate=1000000 sensitive=false
  synonyms={DYNAMIC_BROKER_CONFIG:leader.replication.throttled.rate=100000}
  follower.replication.throttled.rate=1000000 sensitive=false
  synonyms={DYNAMIC_BROKER_CONFIG:follower.replication.throttled.rate=1000000}
Dynamic configs for broker 103 are:
  leader.replication.throttled.rate=1000000 sensitive=false
  synonyms={DYNAMIC_BROKER_CONFIG:leader.replication.throttled.rate=100000}
  follower.replication.throttled.rate=1000000 sensitive=false
  synonyms={DYNAMIC_BROKER_CONFIG:follower.replication.throttled.rate=100000}
Default configs for brokers in the cluster are:
```

10. Verify that throttling is in effect for topic **moving**:

```
$ kafka-configs      --bootstrap-server kafka-1:19092      --describe
--topic moving
Dynamic configs for topic moving are:

  leader.replication.throttled.replicas=0:101,0:102,1:101,1:102,2:101,2
  :102,4:101,4:102 sensitive=false
  synonyms={DYNAMIC_TOPIC_CONFIG:leader.replication.throttled.replicas=
  0:101,0:102,1:101,1:102,2:101,2:102,4:101,4:102}
  follower.replication.throttled.replicas=0:103,1:103,2:103,4:103
  sensitive=false
  synonyms={DYNAMIC_TOPIC_CONFIG:follower.replication.throttled.replica
  s=0:103,1:103,2:103,4:103}
```

11. Monitor the progress of the rebalancing:

```
$ confluent-rebalancer status \
--bootstrap-server kafka-1:19092
Partitions being rebalanced:
Topic moving: 0,1,2,4
```



Your status may indicate different partitions being rebalanced.

12. Increase the throttle limit configuration to 1GBps by rerunning the **confluent-rebalancer** command with the new throttle limit:

```
$ confluent-rebalancer execute \
--bootstrap-server kafka-1:19092 \
--metrics-bootstrap-server kafka-1:19092,kafka-2:29092 \
--throttle 1000000000 \
--verbose
The throttle rate was updated to 1000000000 bytes/sec.
A rebalance is currently in progress for:
Topic moving: 0,1,2,4
```

13. Note the updated throttle limit configuration values.

```

$ kafka-configs \
  --describe \
  --bootstrap-server kafka-1:19092 \
  --entity-type brokers
Dynamic configs for broker 101 are:
  leader.replication.throttled.rate=1000000000 sensitive=false
  synonyms={DYNAMIC_BROKER_CONFIG:leader.replication.throttled.rate=100
  0000000}
  follower.replication.throttled.rate=1000000000 sensitive=false
  synonyms={DYNAMIC_BROKER_CONFIG:follower.replication.throttled.rate=1
  000000000}
Dynamic configs for broker 102 are:
  leader.replication.throttled.rate=1000000000 sensitive=false
  synonyms={DYNAMIC_BROKER_CONFIG:leader.replication.throttled.rate=100
  0000000}
  follower.replication.throttled.rate=1000000000 sensitive=false
  synonyms={DYNAMIC_BROKER_CONFIG:follower.replication.throttled.rate=1
  000000000}
Dynamic configs for broker 103 are:
  leader.replication.throttled.rate=1000000000 sensitive=false
  synonyms={DYNAMIC_BROKER_CONFIG:leader.replication.throttled.rate=100
  0000000}
  follower.replication.throttled.rate=1000000000 sensitive=false
  synonyms={DYNAMIC_BROKER_CONFIG:follower.replication.throttled.rate=1
  000000000}
Default configs for brokers in the cluster are:

```

14. Check the status of the Auto Data Balancer again.

```

$ confluent-rebalancer status \
  --bootstrap-server kafka-1:19092

```

Note that it has completed:

Error: No rebalance is currently in progress. If you have called `status` after a rebalance was started successfully, the rebalance has completed. Run the `execute` command to check if the cluster is balanced.

1. Check to see if topic **moving** is still throttled:

```

$ kafka-configs \
  --bootstrap-server kafka-1:19092 \
  --describe \
  --topic moving

```

Note that all throttles have been removed:

Dynamic configs for topic moving are:

2. Wait up to five minutes. In the Control Center **Broker overview**, look at the metric **Disk**. Notice the red bar is no longer present.
3. View the Topic information to see that the Topic **moving** has its Partitions moved across all three Brokers.

Partitions		Replica placement		Offset Start	End	Size Total Size
Partition id	Status	Leader (broker ID)	Followers (broker IDs)			
0	Available	101	103	0	333333	338MB
1	Available	103	101	0	333334	338MB
2	Available	103	102	0	333334	338MB
3	Available	102	103	0	333333	338MB
4	Available	101	102	0	333333	338MB

4. Return to the terminal running the Consumer. Press **Ctrl+C** to exit the Consumer.

Simulate a Completely Failed Broker

In previous exercises, the Broker failures were such that the Broker could be simply restarted to recover. Now you will simulate a completely failed Broker which requires a replacement system.

1. In a new terminal from the host, observe the logs on Broker 1.

```
$ docker-compose exec kafka-1 ls /var/lib/kafka/data
...
moving-2
moving-3
moving-4
moving-5
recovery-point-offset-checkpoint
replication-offset-checkpoint
...
```

2. Stop and remove Broker 1 (this also removes the data of Broker 1).

```
$ docker-compose stop kafka-1 && \
  docker-compose rm kafka-1 && \
  docker volume rm confluent-admin_data-kafka-1
Stopping kafka-1 ... done

Going to remove kafka-1
Are you sure? [yN] y
Removing kafka-1 ... done

confluent-admin_data-kafka-1
```

3. In the Control Center under **Overview**, verify that only 2 Brokers are available.



You may have to wait up to five minutes for the Broker to disappear.

4. Restart Broker 1:

```
$ docker-compose up -d

control-center is up-to-date
zk-1 is up-to-date
zk-2 is up-to-date
ksql-cli is up-to-date
ksql-server is up-to-date
zk-3 is up-to-date
kafka-2 is up-to-date
schema-registry is up-to-date
kafka-3 is up-to-date
connect is up-to-date
Creating kafka-1 ... done
```

5. In the Control Center under **Overview** verify that all three Brokers are running.



You may have to wait up to five minutes for the previously failed Broker to reappear.

6. Notice that there are under replicated and offline Topic Partitions. After a moment of recovery they will return back to zero though.
7. Look at the logs in `/var/lib/kafka/data`. Verify that Broker 1 has not lost any of the data it had in step 1:

```
$ docker-compose exec kafka-1 ls /var/lib/kafka/data
```



STOP HERE. THIS IS THE END OF THE EXERCISE.

Lab 07 Consumer Groups and Load Balancing

a. Modifying Partitions and Viewing Offsets

In this exercise, you will increase the number of Partitions in a Topic and view offsets in an active Consumer Group.

Prerequisites

1. Please make sure you have prepared your environment by following → [Preparing the Labs](#)
2. Make sure your Kafka cluster is started, **otherwise** execute this command:

```
$ cd ~/confluent-admin  
$ docker-compose up -d
```

Increasing the Number of Partitions in a Topic

1. Create a new Topic called **grow-topic** with six Partitions and three replicas:

```
$ kafka-topics \  
  --bootstrap-server kafka-1:19092 \  
  --create \  
  --topic grow-topic \  
  --partitions 6 \  
  --replication-factor 3
```

2. View the **grow-topic** configuration and replica placement in the cluster:
 - a. In Control Center make sure you select the cluster **Cluster 1**
 - b. Then, on the left side click on **Topics** and then click on the **grow-topic** Topic to see:

Partitions	Partition id	Status	Replica placement	Leader (broker ID)	Followers (broker IDs)	Offset	Start	End	Size	Total Size
	0	Available		101	103, 102		0	0	0B	0B
	1	Available		103	102, 101		0	0	0B	0B
	2	Available		102	101, 103		0	0	0B	0B
	3	Available		101	102, 103		0	0	0B	0B
	4	Available		103	101, 102		0	0	0B	0B
	5	Available		102	103, 101		0	0	0B	0B

3. Start the console Producer for Topic **grow-topic**.

```
$ kafka-console-producer \
--bootstrap-server kafka-1:19092,kafka-2:29092 \
--topic grow-topic
```

At the **>** prompt, type three messages and then press **Ctrl+D** to exit the console Producer:

```
> ksqlDB
> Streaming
> Engine
<Ctrl+D>
```

4. Start the console Consumer for Topic **grow-topic** specifying the **group.id** property:

```
$ kafka-console-consumer \
  --consumer-property group.id=test-consumer-group \
  --from-beginning \
  --topic grow-topic \
  --bootstrap-server kafka-1:19092,kafka-2:29092
```

Streaming
ksqlDB
Engine

Leave this Consumer running during the next step. Think about why it needs to keep running.

- From another terminal window, describe the Consumer Group **test-consumer-group**:

```
$ kafka-consumer-groups \
  --bootstrap-server kafka-1:19092,kafka-2:29092 \
  --group test-consumer-group \
  --describe
```

TOPIC	PARTITION	CURRENT-OFFSET	LOG-END-OFFSET	LAG	CONSUMER-ID
HOST	CLIENT-ID				
grow-topic	1	1	1	0	consumer-1-
a9e... /192.168.0.6	consumer-1				
grow-topic	4	0	0	0	consumer-1-
a9e... /192.168.0.6	consumer-1				
grow-topic	0	1	1	0	consumer-1-
a9e... /192.168.0.6	consumer-1				
grow-topic	3	1	1	0	consumer-1-
a9e... /192.168.0.6	consumer-1				
grow-topic	2	0	0	0	consumer-1-
a9e... /192.168.0.6	consumer-1				
grow-topic	5	0	0	0	consumer-1-
a9e... /192.168.0.6	consumer-1				

- Alter the Topic to increase the number of Partitions to 12:

```
$ kafka-topics \
  --bootstrap-server kafka-1:19092 \
  --alter \
  --topic grow-topic \
  --partitions 12
```



If partitions are increased for a topic that has a key, the partition logic or ordering of the messages will be affected!

7. Wait for a refresh of the metadata, which happens after `metadata.max.age.ms` (default: 5 minutes). Then describe the Consumer Group `test-consumer-group` again. What has changed?

```
$ kafka-consumer-groups \
  --bootstrap-server kafka-1:19092,kafka-2:29092 \
  --group test-consumer-group \
  --describe
```

8. In Control Center, view the `grow-topic` topic configuration from the `Topics` view. You may need to refresh your browser:

Partitions Partition id	Status	Replica placement	
		Broker leader	Broker followers
7	Available	101	103, 102
8	Available	102	101, 103
9	Available	103	101, 102
10	Available	101	102, 103
11	Available	102	103, 101



You need to scroll the list of partitions to see all partitions.

9. In the terminal window where the Consumer is running, press `Ctrl+C` to terminate the process.

Kafka-based Offset Storage

1. Start the console Producer for Topic `new-topic`:

```
$ kafka-console-producer \
  --bootstrap-server kafka-1:19092,kafka-2:29092 \
  --topic new-topic
```

At the `>` prompt, type some messages into the console. After the first message, you will see a warning message because the Topic `new-topic` did not exist prior to producing to the Topic. Press `Ctrl+D` to return to the command line when you are done:

```
> I  
[2018-08-02 14:01:26,083] WARN [Producer clientId=console-producer]  
Error while fetching metadata with correlation id 2 : {new-  
topic=LEADER_NOT_AVAILABLE} (org.apache.kafka.clients.NetworkClient)  
> Love  
> Kafka  
<Ctrl+D>
```

2. Start the console Consumer for Topic `--consumer_offsets` to view the offsets. Leave this running for the duration of the exercise:

```
$ kafka-console-consumer \  
  --topic __consumer_offsets \  
  --bootstrap-server kafka-1:19092,kafka-2:29092 \  
  --formatter  
"kafka.coordinator.group.GroupMetadataManager\$OffsetsMessageFormatte  
r" \  
  | grep new-topic
```

3. Start the console Consumer in a new terminal window for Topic `new-topic` in a Consumer Group called `new-group`:

```
$ kafka-console-consumer \  
  --from-beginning \  
  --topic new-topic \  
  --group new-group \  
  --bootstrap-server kafka-1:19092,kafka-2:29092  
I  
Love  
Kafka
```

4. Press `Ctrl+C` to terminate the Consumer once your messages have been displayed.
 - a. When reading from the Topic `new-topic`, did you see the messages you typed earlier?
 - b. In the other terminal window where you were reading from the Topic `--consumer_offsets`, did you see any `OffsetMetadata` specifically for the Topic `new-topic` such as shown below?

```
[new-group,new-topic,0]::OffsetAndMetadata(offset=3,  
leaderEpoch=Optional[0], metadata=, commitTimestamp=1563289258184,  
expireTimestamp=None)  
[new-group,new-topic,0]::OffsetAndMetadata(offset=3,  
leaderEpoch=Optional[0], metadata=, commitTimestamp=1563289263184,  
expireTimestamp=None)  
...
```

5. If they are currently running, terminate the Producer and Consumers with **Ctrl+C**.



STOP HERE. THIS IS THE END OF THE EXERCISE.

Lab 08 Optimizing Kafka's Performance

a. Exploring Producer Performance

Prerequisites

1. Please make sure you have prepared your environment by following → [Preparing the Labs](#)
2. Make sure your Kafka cluster is started, **otherwise** execute this command:

```
$ cd ~/confluent-admin  
$ docker-compose up -d
```

Tuning Producer Performance

1. Create a new Topic called **performance** with six Partitions and three replicas.

```
$ kafka-topics \  
  --bootstrap-server kafka-1:19092 \  
  --create \  
  --topic performance \  
  --partitions 6 \  
  --replication-factor 3
```

2. Run Producer performance tests to compare **acks=all**, **acks=1**, and **acks=0**

Use the following command line and replace **<VARIABLE_HERE>** with the specified options in the table below:

```
$ kafka-producer-perf-test \
  --topic performance \
  --num-records 1000000 \
  --record-size 100 \
  --throughput 10000000 \
  --producer-props \
  bootstrap.servers=kafka-1:19092,kafka-2:29092 \
  <VARIABLE_HERE>
```

Let each test run until it terminates and provides a performance summary. Record the throughput and latency results in a table. You will be appending to this table in upcoming questions.

Variables	Throughput (MB/sec)	Latency (ms avg)
acks=1		
acks=all		
acks=0		

3. Questions:

- Do you get better throughput with **acks=0**, **acks=1**, or **acks=all**? What is the percentage difference in performance?
 - Do you get better latency with **acks=0**, **acks=1**, or **acks=all**? Why?
4. You will investigate the effects of tuning **batch.size** and **linger.ms** by investigating **one** question. Try values of **batch.size** between 0 and 1,000,000 and values of **linger.ms** between 0 and 3,000. Hold **acks=all** constant. **Choose one question to investigate.** Consider dividing work amongst peers who are investigating the same question. Record your results in a table.
- What is the maximum throughput, no matter the latency?
 - What is the minimum latency, no matter the throughput?
 - What is the best balance of throughput and latency? (and defend your decision)
 - Given a batch size of 100,000 Bytes, what linger gives best performance? What do you notice? What do you wonder?
 - Given a linger of 500 ms, what batch size gives best performance? What do you notice? What do you wonder?

Here is an example of a test you might run:

```
$ kafka-producer-perf-test \
--topic performance \
--num-records 1000000 \
--record-size 100 \
--throughput 10000000 \
--producer-props \
bootstrap.servers=kafka-1:19092,kafka-2:29092 \
acks=all \
batch.size=400000 \
linger.ms=500
```



STOP HERE. THIS IS THE END OF THE EXERCISE.

b. Performance Tuning

In this exercise, you will observe Kafka performance and use some of Kafka's settings to monitor and optimize Consumers.

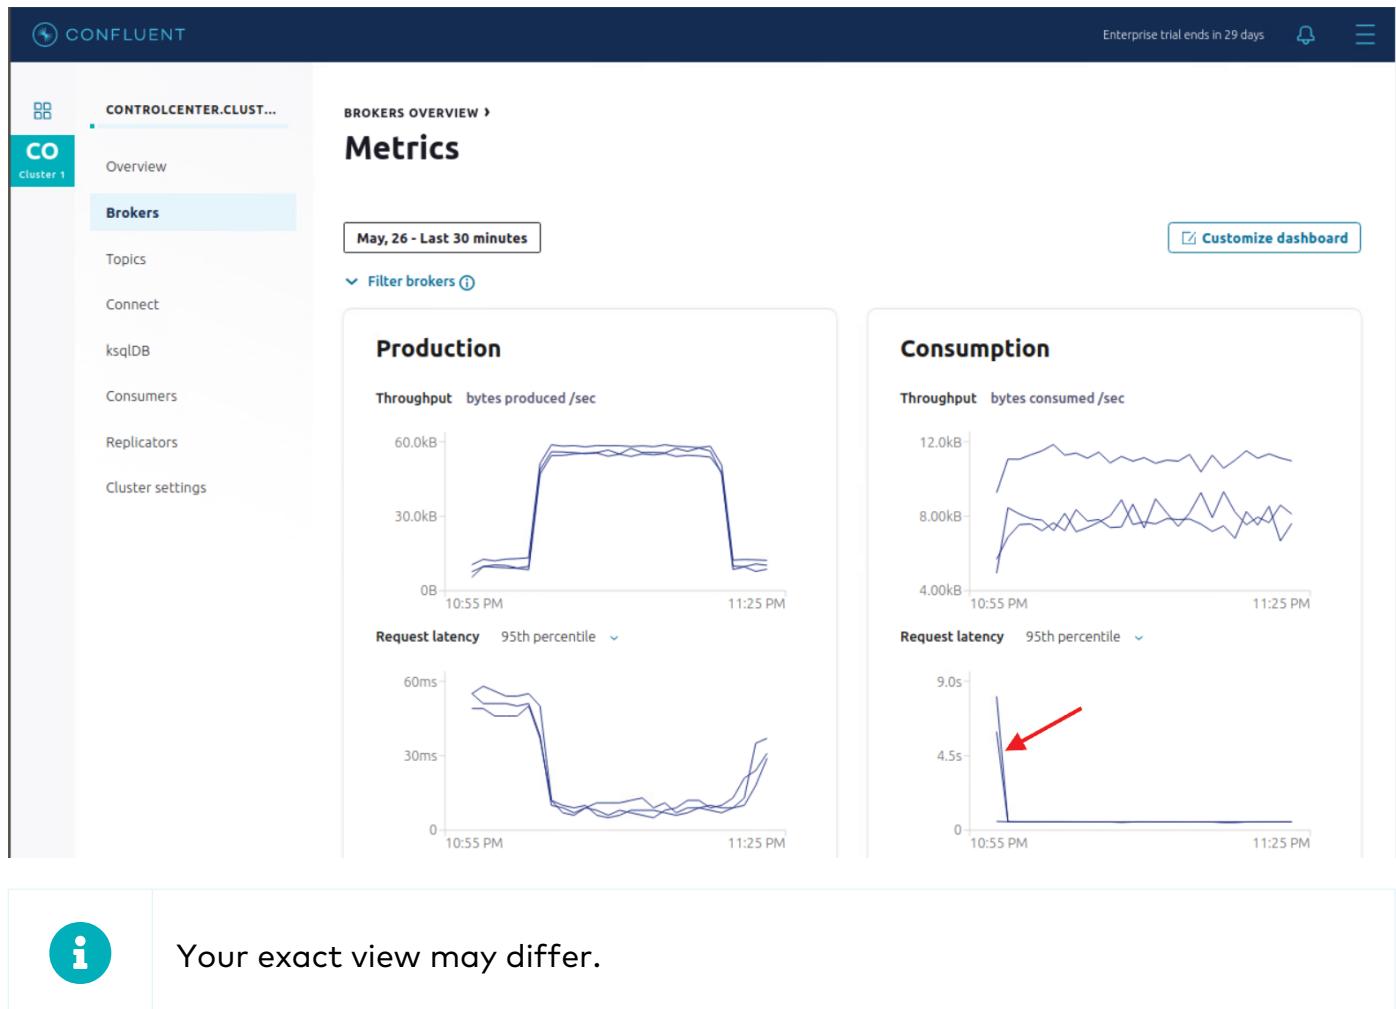
Prerequisites

1. Please make sure you have prepared your environment by following → [Preparing the Labs](#)
2. Make sure your Kafka cluster is started, **otherwise** execute this command:

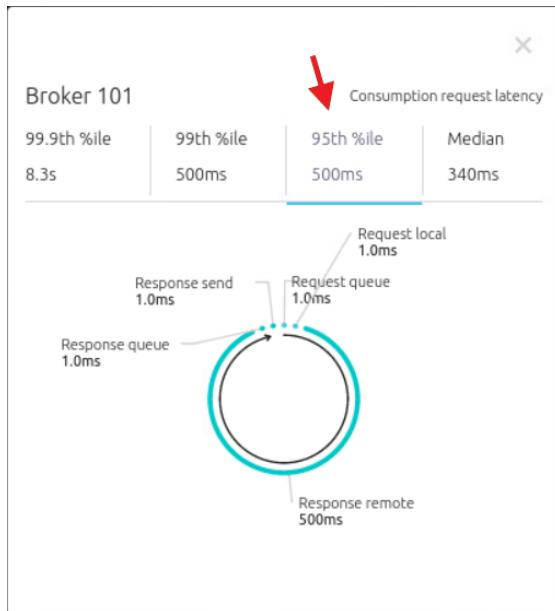
```
$ cd ~/confluent-admin
$ docker-compose up -d
```

Observing Replica Fetch Times

1. Connect to Control Center. If necessary, open a browser tab to the URL
<http://localhost:9021>.
2. In the Control Center, go to **Cluster 1** → **Brokers** → **Consumption**. Click on one of the lines in the consumption request latency line graph to open the breakdown of times in the request latency lifecycle.



3. In the fetch request latency breakdown, select the **95th %ile** view.



Remember that percentiles don't add associatively, so the numbers around the circle won't generally add up to the overall request time percentile.

4. Observe that most of the fetch request time is in **Response remote** time waiting for the **replica.fetch.wait.max.ms** timeout. When Producers are not writing records, the fetch request latency will go up to **replica.fetch.wait.max.ms** (default 500ms).
5. Create a new Topic called **fetch-request** with six Partitions and three replicas:

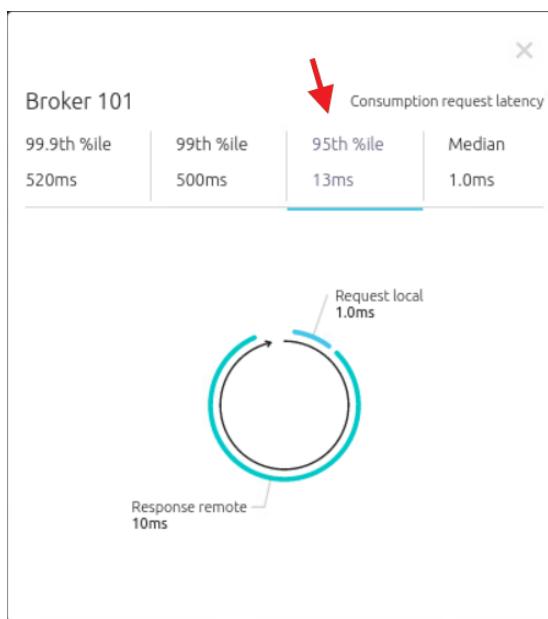
```
$ kafka-topics \
--bootstrap-server kafka-1:19092 \
--create \
--topic fetch-request \
--partitions 6 \
--replication-factor 3
```

6. Produce some data to send to the Topic **fetch-request** and let it run:

```
$ kafka-producer-perf-test \
--topic fetch-request \
--num-records 1000000 --record-size 100 --throughput 1000 \
--producer-props bootstrap.servers=kafka-1:19092,kafka-2:29092

4988 records sent, 997.6 records/sec (0.10 MB/sec), 19.3 ms avg
latency, 278.0 max latency.
5004 records sent, 999.2 records/sec (0.10 MB/sec), 10.5 ms avg
latency, 83.0 max latency.
5028 records sent, 1005.4 records/sec (0.10 MB/sec), 12.6 ms avg
latency, 123.0 max latency.
...
```

- Let the producer run for five minutes. In Control Center's fetch request latency breakdown, observe that the **response remote time** at the **95th %ile** drops to a few milliseconds:



- If it is currently running, terminate the Producer by pressing **Ctrl+C**.

Tune Consumers to Decrease Broker CPU Load

- Create a new Topic called **i-love-logs** with one Partition and one replica.

```
$ kafka-topics \
--bootstrap-server kafka-1:19092 \
--create \
--topic i-love-logs \
--replica-assignment 101
```



We are using the `--replica-assignment` parameter to insure that the partition for all students is located on the `kafka-1` broker which is assigned broker id `101`.

2. Produce some data to send to the Topic `i-love-logs`. Leave this process running until instructed to end it.

The following command line sends 10 million messages, 100 bytes in size, at a rate of 1000 messages/sec. The `NS` environment variable defines interceptors that enable stream monitoring in Control Center.

```
$ NS=io.confluent.monitoring.clients.interceptor && \
kafka-producer-perf-test \
--topic i-love-logs \
--num-records 10000000 \
--record-size 100 \
--throughput 1000 \
--producer-props \
    bootstrap.servers=kafka-1:19092,kafka-2:29092 \
    interceptor.classes=${NS}.MonitoringProducerInterceptor
```

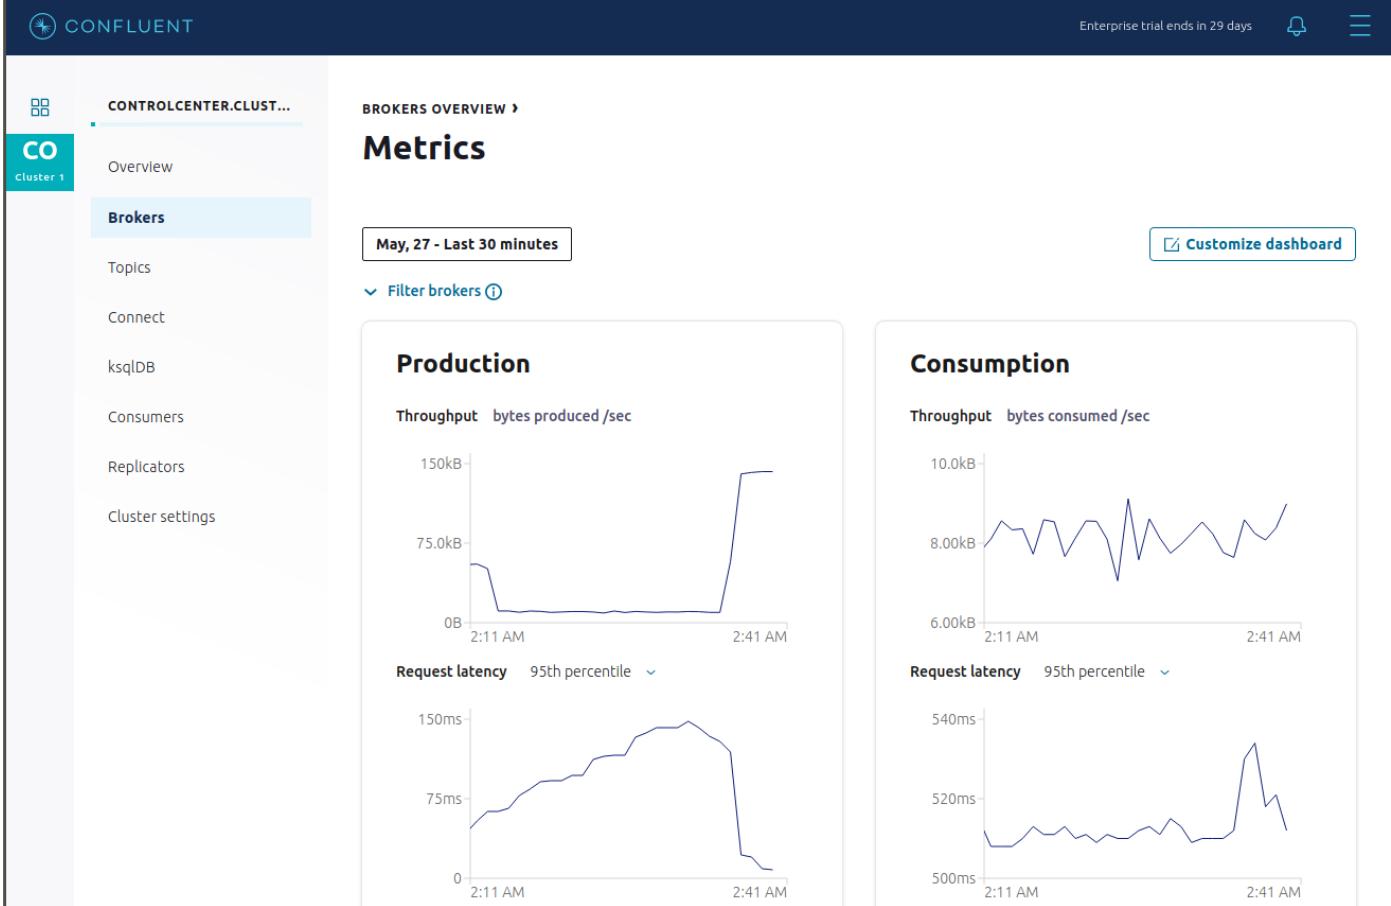
5001 records sent, 1000.0 records/sec (0.10 MB/sec), 7.9 ms avg latency, 371.0 max latency.

5004 records sent, 1000.6 records/sec (0.10 MB/sec), 1.4 ms avg latency, 31.0 max latency.

5000 records sent, 1000.0 records/sec (0.10 MB/sec), 1.2 ms avg latency, 7.0 max latency.

...

3. Wait for a few minutes. Go to `Cluster 1 → Brokers → Production` in Control Center. If necessary, adjust the timescale of the data to the last 30 minutes.
4. Select `broker.id 101` in the `Filter brokers` list.
5. Observe changes that occur in the `Throughput` and `Request latency`. After five minutes, the graphs should look similar to this:



6. Open a new terminal and navigate to the `~/confluent-admin` directory:

```
$ cd ~/confluent-admin
```

7. Create a consumer properties file called `data/consumer.properties` that will enable Control Center to monitor Consumer performance of Consumer Group `cg`:

```
$ PACKAGE=io.confluent.monitoring.clients.interceptor && \
echo "interceptor.classes=${PACKAGE}.MonitoringConsumerInterceptor" \
> data/consumer.properties
```

8. Consume some data from the Topic `i-love-logs` with Consumer Group `cg`. Leave this process running until instructed to end it.

```

$ kafka-consumer-perf-test \
    --bootstrap-server kafka-1:19092,kafka-2:29092 \
    --topic i-love-logs \
    --group cg \
    --messages 10000000 \
    --threads 1 \
    --show-detailed-stats \
    --reporting-interval 5000 \
    --consumer.config data/consumer.properties

time, threadId, data.consumed.in.MB, MB.sec, data.consumed.in.nMsg,
nMsg.sec, rebalance.time.ms, fetch.time.ms, fetch.MB.sec,
fetch.nMsg.sec
2018-08-06 12:24:33:271, 0, 26.2823, 5.2554, 275590, 55106.9786,
3072, 1929, 13.6248, 142866.7703
2018-08-06 12:24:38:271, 0, 118.8769, 18.5189, 1246515, 194185.0000,
0, 5000, 18.5189, 194185.0000
2018-08-06 12:24:43:274, 0, 218.6112, 19.9349, 2292305, 209032.5805,
0, 5003, 19.9349, 209032.5805
...

```

- In a new terminal, observe the CPU load for the **java** process on Broker **kafka-1**:

```

$ cd ~/confluent-admin
$ docker-compose exec kafka-1 top -n10
top - 14:09:16 up 13 min,  0 users,  load average: 5.75, 7.06, 3.90
Tasks:  2 total,   1 running,   1 sleeping,   0 stopped,   0 zombie
%Cpu(s): 65.8 us, 24.6 sy,  0.0 ni,  5.7 id,  0.3 wa,  0.0 hi,  3.5
si,  0.0 st
MiB Mem : 9992.6 total,   2216.4 free,   5589.9 used,   2186.3
buff/cache
MiB Swap:      0.0 total,      0.0 free,      0.0 used.  4043.0
avail Mem

      PID USER      PR  NI      VIRT      RES      SHR S %CPU %MEM     TIME+
COMMAND
      1 appuser    20    0  8398896 486448  21392 S 39.3  4.8  1:43.24
java
     169 appuser   20    0    50284   4104   3516 R  0.0  0.0  0:00.39
top

```

Observe CPU utilization until **top** exits. In the output above, the java process CPU load is 39.3%.

- Go to the terminal with the Consumer Group called **cg** and kill it with **Ctrl+C**.
- Add consumer setting **fetch.min.bytes=10485760** to **data/consumer.properties**:

```
$ PACKAGE=io.confluent.monitoring.clients.interceptor && \
echo "interceptor.classes=${PACKAGE}.MonitoringConsumerInterceptor" \
> data/consumer.properties && \
echo "fetch.min.bytes=10485760" >> data/consumer.properties
```

This configuration tells the Broker to wait for larger amounts of data to accumulate before responding to the Consumer. This generally improves throughput and reduces load on the Broker but can increase latency.

12. Continue consuming data from for the Topic **i-love-logs** with Consumer Group **cg** using the updated **consumer.properties**:

```
$ kafka-consumer-perf-test \
    --bootstrap-server kafka-1:19092,kafka-2:29092 \
    --topic i-love-logs \
    --group cg \
    --messages 10000000 \
    --threads 1 \
    --show-detailed-stats \
    --reporting-interval 5000 \
    --consumer.config data/consumer.properties
```

13. In a different terminal from the one running the **cg** Consumer Group, observe the CPU load again for the **java** process on Broker **kafka-1**:

```
$ cd ~/confluent-admin
$ docker-compose exec kafka-1 top -n10
Tasks:  2 total,  1 running,  1 sleeping,  0 stopped,  0 zombie
%Cpu(s): 20.9 us,  8.2 sy,  0.0 ni, 69.5 id,  0.0 wa,  0.0 hi,  1.4
si,  0.0 st
MiB Mem : 9992.6 total,   2091.8 free,   5641.2 used,   2259.6
buff/cache
MiB Swap:      0.0 total,      0.0 free,      0.0 used. 3990.9
avail Mem

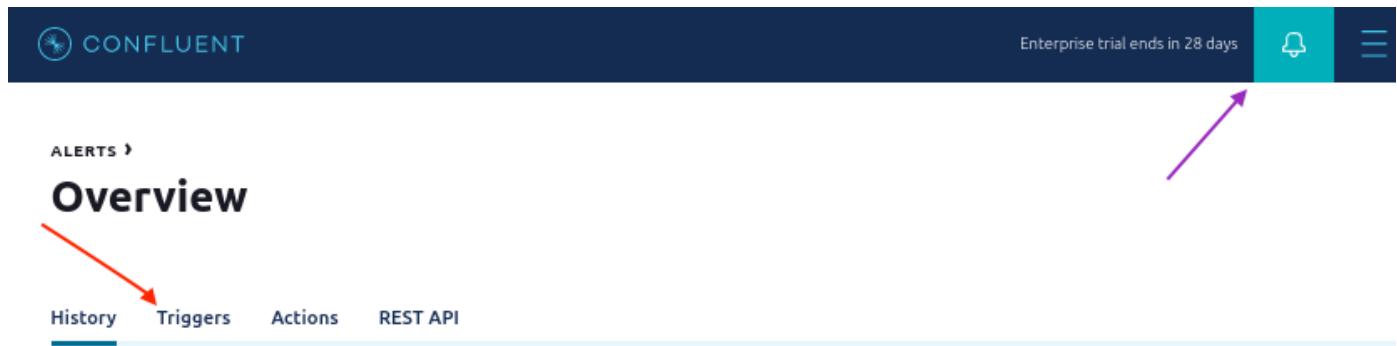
      PID USER      PR  NI      VIRT      RES      SHR S %CPU %MEM     TIME+
COMMAND
      1 appuser    20    0  8398896  501376  21660 S 12.3  4.9  2:35.39
java
  174 appuser    20    0    50284    4144    3560 R  0.0  0.0  0:00.31
top
```

Indeed, increasing **fetch.min.bytes** reduces the CPU load of the Broker.

Defining Over Consumption Trigger and Action

For this exercise, we need to define a trigger and corresponding action in Control Center:

1. Open Control Center
2. Click the bell icon to get to **Alerts**:



3. Select the **Triggers** tab
4. Click the button **New trigger**
5. Complete the dialog using the following settings:

Field	Value
Trigger Name	Over Consumption
Component type	Consumer group
Consumer group name	cg ...
Metric	Consumption difference
Buffer (seconds)	120
Condition	Greater than
Value	0

New trigger

General

Trigger name* _____
Over Consumption

Components

Component type* _____
Consumer group

Consumer group name* _____
cg [Ldd0yZgFR7aN4kkkjdzRnw]

Criteria

Metric* _____
Consumption difference

Buffer (seconds)* _____
120

Condition* _____
Greater than

Value* _____
0

Save

Cancel

6. Click **Save**

7. In the confirmation box, when asked to "Set actions to perform when your trigger goes off", select **Create an action**

8. Complete the dialog using the following settings:

Field	Value
Action Name	Over Consumption
Triggers	Over Consumption
Actions	Send email
Subject	Over Consumption
Max send rate	12 Per hour
Recipient email address	anyone@anywhere.com

New action

General

Action Name*

Over Consumption

Enabled

Triggers

Triggers*

Over Consumption x

Actions

Action*

Send email



Subject*

Over Consumption

Max send rate*

12

Per hour



Recipient email address*

anyone@anywhere.com

Save

Cancel

9. Click **Save**

Simulating Over Consumption

1. Go to the terminal with the Consumer Group called **cg** and kill it with **Ctrl+C**.
2. Identify the current offsets of Consumer Group **cg**:

```
$ kafka-consumer-groups \
  --bootstrap-server kafka-1:19092,kafka-2:29092 \
  --group cg \
  --describe

Consumer group 'cg' has no active members.

GROUP          TOPIC      PARTITION  CURRENT-OFFSET  LOG-END-
OFFSET        LAG        CONSUMER-ID    HOST        CLIENT-ID
cg            i-love-logs   0          671615       689170
17555          -          -          -           -
```

3. Reset the offsets of Consumer Group **cg** offset for Partition **0** of Topic **i-log-logs** to **<current offset> - 100**. Using the previous step as an example, this would be, **671615 - 100 = 671515**:

```
$ kafka-consumer-groups \
  --bootstrap-server kafka-1:19092,kafka-2:29092 \
  --group cg \
  --reset-offsets \
  --to-offset <value> \
  --topic i-love-logs:0 \
  --execute

GROUP          TOPIC
PARTITION  NEW-OFFSET
cg            i-love-logs
671515          0
```

4. Continue consuming data from for the Topic **i-love-logs** with Consumer Group **cg**. Since the offsets have been reset to **<current offset> - 100**, the Consumer Group will consume **100** messages from the **i-love-logs** Topic a second time, thereby processing these messages again. This is a good simulation of Consumer Group over-consumption:

```
$ kafka-consumer-perf-test \
    --bootstrap-server kafka-1:19092,kafka-2:29092 \
    --topic i-love-logs \
    --group cg \
    --messages 10000000 \
    --threads 1 \
    --show-detailed-stats \
    --reporting-interval 5000 \
    --consumer.config data/consumer.properties
```

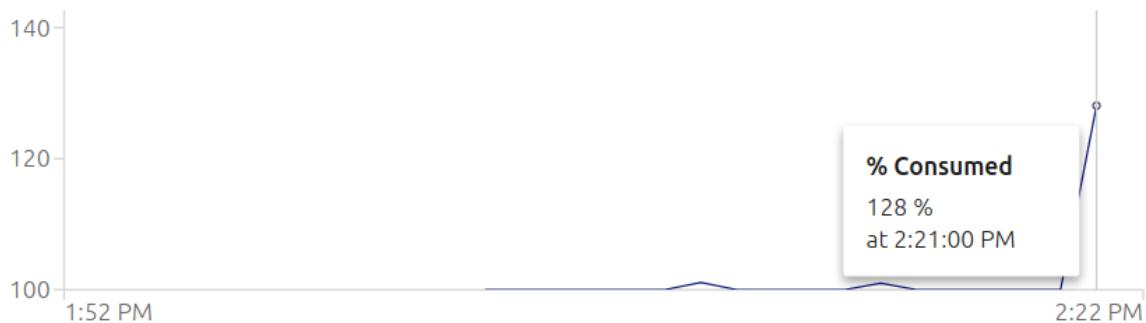
5. Observe the cluster in Control Center for a couple minutes. You may need to adjust your time window.
 - a. Go to **Cluster 1** → **Consumers** → **cg** → **Consumption** and observe the Consumer Group **cg** for a few minutes:

cg

Consumer lag Consumption

Jun, 1 - Last 30 minutes

% messages consumed



End-to-end latency average (ms) ▾



Observe that:

- The **% messages consumed** is over 100%

- Click the bell icon to go back to the **Alerts Overview**. View the alert history and notice that the over-consumption event triggered an alert. It is critical to monitor applications for message consumption to know how your applications are behaving. Over-consumption may happen intentionally, or it may happen unintentionally, for example if an application crashes before committing processed messages.

Overview

History Triggers Actions REST API

Time	Trigger	Component	Action Count
a minute ago	Over Consumption	cg	1

7. End all running Producers and Consumers by going to the respective terminal window and pressing **Ctrl+C**.



STOP HERE. THIS IS THE END OF THE EXERCISE.

Lab 09 Kafka Security

a. Securing the Kafka Cluster

In this exercise, you will configure one-way SSL authentication, two-way SSL authentication, and explore the SSL performance impact.

Prerequisites

1. Please make sure you have prepared your environment by following → [Preparing the Labs](#)
2. If a cluster is currently running, use `docker-compose down` to shut it down.

```
$ cd ~/confluent-admin  
$ docker-compose down -v
```

3. Change directory to the `secure-cluster` folder:

```
$ cd ~/confluent-admin/secure-cluster
```

4. Start the secure cluster:

```
$ docker-compose up -d
```



The broker containers have been configured to wait until their respective `kafka-n-creds` directory has been populated with the required certs. Once the certs are present, the broker startup continues. These certs are being created in the next step of the exercise.

Generating Certificates

Next, we generate all the necessary certificates and credentials that are used to create a secure, 3-Broker Kafka cluster:

1. Generate a signed certificate, keystores, and truststores by running a script:

```
$ ./certs-create.sh
```

2. Use VScode or other editor to inspect the files the script created:

```
$ code .
```

The **ca.key** file is the private key that the Certificate Authority uses to sign the certificate **ca.crt**. The script used this certificate to create the Broker keystore and the client truststore. The Broker's keystore is what the Broker uses to send its certificate to the client. The client truststore is what the client uses to check whether a Broker's certificate ought to be trusted—if the checksum of the received certificate is unexpected, then someone tried to alter the certificate and thus the certificate is untrustworthy.

Later during mutual SSL, the client will also need to authenticate with the Broker. Therefore, the script also created a keystore for the client and a truststore for the Broker.

Enabling SSL on the Brokers

1. Open the file **~/confluent-admin/secure-cluster/docker-compose.yml** in your editor and note the following environment variables (in addition to the existing ones) that have been added to each of the three Brokers (**kafka-1**, **kafka-2** and **kafka-3**), e.g. for **kafka-1** they are:

```
KAFKA_ADVERTISED_LISTENERS: SSL://kafka-1:19093,PLAINTEXT://kafka-1:19092,DOCKER://kafka-1:9092
KAFKA_LISTENER_SECURITY_PROTOCOL_MAP:
  PLAINTEXT:PLAINTEXT,SSL:SSL,DOCKER:PLAINTEXT
KAFKA_SSL_KEYSTORE_FILENAME: kafka.kafka-1.keystore.jks
KAFKA_SSL_KEYSTORE_CREDENTIALS: kafka-1_keystore_creds
KAFKA_SSL_KEY_CREDENTIALS: kafka-1_sslkey_creds
KAFKA_SSL_TRUSTSTORE_FILENAME: kafka.kafka-1.truststore.jks
KAFKA_SSL_TRUSTSTORE_CREDENTIALS: kafka-1_truststore_creds
KAFKA_SSL_ENDPOINT_IDENTIFICATION_ALGORITHM: "HTTPS"
```

The variable names shown above might not look familiar at first. Your trainer has shown you variable names in the configuration files for Brokers such as **ssl.keystore.filename** or **ssl.truststore.credentials**. A startup script inside each container converts the environment variable names by a standard algorithm into the Kafka specific variable names and adds them to the respective configuration file. Upper case is converted to lowercase and underscores are converted to periods. Furthermore each component (Kafka, REST Proxy, Kafka Connect, etc.) has a well defined prefix that is removed from the Docker variable. In the above case it is **KAFKA_**.

2. Each of the three Brokers maps the credential folder into **/etc/kafka/secrets**. For example, **kafka-1** has:

```
volumes:  
  ...  
  - $PWD/kafka-1-creds:/etc/kafka/secrets
```

3. Please answer the following questions. Feel free to discuss with peers:
 - a. Why is **KAFKA_ADVERTISED_LISTENERS** configured for both **SSL** and **PLAINTEXT**?
 - b. What would happen if **KAFKA_ADVERTISED_LISTENERS** were configured for just **SSL**?
 - c. What does **KAFKA_SSL_ENDPOINT_IDENTIFICATION_ALGORITHM** do?

Verifying SSL is Working

1. From a terminal window verify that the keystore and truststore of each Broker are setup properly. The **openssl** command below should return a key and certificate:

```
$ openssl s_client -connect kafka-1:19093 -tls1_2
CONNECTED(00000005)
depth=1 CN = ca1.test.confluent.io, OU = TEST, O = CONFLUENT, L =
PaloAlto, C = US
verify error:num=19:self signed certificate in certificate chain
---
Certificate chain
  0 s:C = US, ST = Ca, L = PaloAlto, O = CONFLUENT, OU = TEST, CN =
kafka-1
    i:CN = ca1.test.confluent.io, OU = TEST, O = CONFLUENT, L =
PaloAlto, C = US
  1 s:CN = ca1.test.confluent.io, OU = TEST, O = CONFLUENT, L =
PaloAlto, C = US
    i:CN = ca1.test.confluent.io, OU = TEST, O = CONFLUENT, L =
PaloAlto, C = US
---
Server certificate
-----BEGIN CERTIFICATE-----
MIIDTDCCAjQCFALRToiRQJ+ipSaIZ1efStGSFbUVMA0GCSqGSIb3DQEBCwUAMGMx
HjAcBgNVBAMMFWNhMS50ZXN0LmNvbmcZsdWVudC5pbzENMASGA1UECwwEVEVTVD
MBAGA1UECgwJQ090RkxVRU5UMREwDwYDVQQHDAhQYWxvQWx0bzELMAkGA1UEBhMC
...
...
```



The error in the 3rd line of output is expected since there are self-signed certificates during the certificate creation process. Press **Ctrl+C** to stop the above command.

2. Repeat the same for the other two Brokers:

```
$ openssl s_client -connect kafka-2:29093 -tls1_2
$ openssl s_client -connect kafka-3:39093 -tls1_2
```

If you see a **Connection refused** error, then the Broker is not properly set up for SSL or may need to be restarted with **docker-compose restart <broker>**. We simulate this error here by providing the wrong port:



```
$ openssl s_client -connect kafka-1:9094 -tls1_2  
140067608003008:error:0200206F:system  
library:connect:Connection  
refused:../crypto/bio/b_sock2.c:110:  
140067608003008:error:2008A067:BIO  
routines:BIO_connect:connect  
error:../crypto/bio/b_sock2.c:111:  
connect:errno=111
```

3. Let's try to produce messages to the cluster via the SSL port.

a. Create a new Topic called **ssl-topic** with one Partition and three replicas:

```
$ kafka-topics \  
--bootstrap-server kafka-1:19092 \  
--create \  
--topic ssl-topic \  
--partitions 1 \  
--replication-factor 3
```



The Admin Client is using the PLAINTEXT listener at port 19092 to create the Topic.

b. Start the console Producer for Topic **ssl-topic**. Here the client is connecting to the SSL port **19093** rather than the PLAINTEXT port **19092**:

```
$ kafka-console-producer \  
--bootstrap-server kafka-1:19093,kafka-2:29093 \  
--topic ssl-topic  
>[2020-10-29 18:30:25,757] WARN [Producer clientId=console-producer] Bootstrap broker kafka-1:19093 (id: -1 rack: null) disconnected (org.apache.kafka.clients.NetworkClient)  
[2020-10-29 18:30:26,066] WARN [Producer clientId=console-producer] Bootstrap broker kafka-2:29093 (id: -2 rack: null) disconnected (org.apache.kafka.clients.NetworkClient)
```

c. Press **Ctrl+C** to stop the producer.

This leads to a connection error. Why?



When clients communicate with Kafka brokers using SSL, we must provide the required configuration settings.

4. We will now configure the Kafka clients to use the truststore.

- a. First, let's examine the **client-ssl.properties** file that we need to specify when we run the **kafka-console-producer** and **kafka-console-consumer** clients when connecting to the broker using SSL:

```
$ cat client-creds/client_ssl.properties
security.protocol=SSL
ssl.truststore.location=client-creds/kafka.client.truststore.jks
ssl.truststore.password=confluent
```

- b. Try to produce messages again using the **client_ssl.properties** file:

```
$ kafka-console-producer \
  --bootstrap-server kafka-1:19093,kafka-2:29093 \
  --topic ssl-topic \
  --producer.config client-creds/client_ssl.properties
```

- c. At the > prompt, type "Security is good" and press Enter. Add a couple more messages and then press **Ctrl+D** to exit the console Producer.

```
> Security is good
> Certificates
> Keys
<Ctrl+D>
```

- d. Start the console Consumer for Topic **ssl-topic** using the **client_ssl.properties** file. You should see the messages you typed above. Press **Ctrl+C** once your messages have been displayed.

```
$ kafka-console-consumer \
  --consumer.config client-creds/client_ssl.properties \
  --from-beginning \
  --topic ssl-topic \
  --bootstrap-server kafka-1:19093,kafka-2:29093
```

Security is good
Certificates
Keys

Press **Ctrl+C** to stop the Consumer.

Enabling Mutual SSL Authentication

1. In the **docker-compose.yml** file in folder **~/confluent-admin/secure-cluster**, uncomment the following environment variable for each Broker:

```
KAFKA_SSL_CLIENT_AUTH: "required"
```

2. Restart all 3 Brokers by executing:

```
$ cd ~/confluent-admin/secure-cluster
$ docker-compose up -d

zk-1 is up-to-date
control-center is up-to-date
schema-registry is up-to-date
zk-3 is up-to-date
Recreating kafka-2 ...
secure-cluster_base_1 is up-to-date
Recreating kafka-2 ... done
Recreating kafka-1 ... done
Recreating kafka-3 ... done
```



Docker will realize that the definition of the 3 Broker definitions have changed and reschedule the respective container.

3. Try to run the **kafka-console-consumer**:

```
$ kafka-console-consumer \
  --consumer.config client-creds/client_ssl.properties \
  --from-beginning \
  --topic ssl-topic \
  --bootstrap-server kafka-1:19093,kafka-2:29093
[2020-10-29 18:47:01,995] ERROR [Consumer clientId=consumer-console-consumer-86926-1, groupId=console-consumer-86926] Connection to node -2 (kafka-2/127.0.0.1:29093) failed authentication due to: Failed to process post-handshake messages
(org.apache.kafka.clients.NetworkClient)
...
...
```

This fails with a Java IO Exception. Why?

Previously, the Broker had to authenticate with the Clients. Now that we configured the Brokers to require mutual SSL, the Clients must also authenticate with Brokers.



The `certs-create.sh` script already created a keystore in the `client-creds` folder and corresponding truststore in each `kafka-{1|2|3}-creds` folders. We need to add the related keystore configuration settings to `client_ssl.properties`.

4. Add the required lines to `client_ssl.properties`:

```
$ echo "ssl.keystore.location=client-creds/kafka.client.keystore.jks" \
>> client-creds/client_ssl.properties && \
echo "ssl.keystore.password=confluent" >> client-creds/client_ssl.properties
```

a. Let's confirm `client-ssl.properties` now contains all of the required settings:

```
$ cat client-creds/client_ssl.properties
security.protocol=SSL
ssl.truststore.location=client-creds/kafka.client.truststore.jks
ssl.truststore.password=confluent
ssl.keystore.location=client-creds/kafka.client.keystore.jks
ssl.keystore.password=confluent
```

5. Start the console Consumer for Topic `ssl-topic` again, this time passing in the updated `client_ssl.properties` file. Now it should succeed and you should see the messages

you typed earlier:

```
$ cd ~/confluent-admin/secure-cluster
$ kafka-console-consumer \
  --consumer.config client-creds/client_ssl.properties \
  --from-beginning \
  --topic ssl-topic \
  --bootstrap-server kafka-1:19093,kafka-2:29093

Security is good
Certificates
Keys
```

Press **Ctrl+C** to quit the Consumer once your messages have been displayed.

SSL Performance Impact

In this section, you will run two performance tests and compare the results:

- Connecting to the cluster with no SSL via the configured **PLAINTEXT** port (19092/29092)
 - Connecting to the cluster with SSL via the configured **SSL** port (19093/29093)
1. Create a new Topic called **no-ssl-topic** with one Partition and three replicas.

```
$ kafka-topics \
  --bootstrap-server kafka-2:29092 \
  --create \
  --topic no-ssl-topic \
  --partitions 1 \
  --replication-factor 3
```

2. Produce a high rate of data to Topic **no-ssl-topic** without SSL. Notice that the Broker port numbers used below are **19092/29092**, the PLAINTEXT port.

```
$ kafka-producer-perf-test \
  --topic no-ssl-topic \
  --num-records 400000 \
  --record-size 1000 \
  --throughput 1000000 \
  --producer-props bootstrap.servers=kafka-1:19092,kafka-2:29092
56833 records sent, 11366.6 records/sec (10.84 MB/sec), 1930.7 ms avg
latency, 2817.0 ms max latency.
67152 records sent, 13427.7 records/sec (12.81 MB/sec), 2373.1 ms avg
latency, 2607.0 ms max latency.
138640 records sent, 27728.0 records/sec (26.44 MB/sec), 1374.8 ms
avg latency, 2488.0 ms max latency.
400000 records sent, 21917.808219 records/sec (20.90 MB/sec), 1416.54
ms avg latency, 2817.00 ms max latency, 1072 ms 50th, 2523 ms 95th,
2688 ms 99th, 2802 ms 99.9th.
```



Performance results will vary.

- a. What is the average throughput? Look in the last line for the value associated with **MB/sec**.
- b. What is the average latency? Look in the last line for the value associated with **ms avg latency**.
3. Now produce a high rate of data to Topic **ssl-topic** with SSL. Notice that you now connect to Broker port number **9093**, the SSL port.

```
$ kafka-producer-perf-test \
  --topic ssl-topic \
  --num-records 400000 \
  --record-size 1000 \
  --throughput 1000000 \
  --producer-props bootstrap.servers=kafka-1:19093,kafka-2:29093 \
  --producer.config client-creds/client_ssl.properties
35537 records sent, 7107.4 records/sec (6.78 MB/sec), 2440.7 ms avg
latency, 3680.0 ms max latency.
63904 records sent, 12780.8 records/sec (12.19 MB/sec), 2711.9 ms avg
latency, 3182.0 ms max latency.
110768 records sent, 22153.6 records/sec (21.13 MB/sec), 1662.3 ms
avg latency, 2517.0 ms max latency.
179744 records sent, 35948.8 records/sec (34.28 MB/sec), 937.6 ms avg
latency, 1227.0 ms max latency.
400000 records sent, 19688.915141 records/sec (18.78 MB/sec), 1555.90
ms avg latency, 3680.00 ms max latency, 1216 ms 50th, 3099 ms 95th,
3442 ms 99th, 3658 ms 99.9th.
```



Performance results will vary.

- a. What is the average throughput? Look in the last line for the value associated with **MB/sec**. Is this higher or lower than without SSL?
- b. What is the average latency? Look for the value associated with **ms avg latency**. Is this higher or lower than without SSL?

Cleanup

1. Execute the following command to completely cleanup your environment:

```
$ cd ~/confluent-admin/secure-cluster  
$ docker-compose down -v
```



STOP HERE. THIS IS THE END OF THE EXERCISE.

Lab 10 Data Pipelines with Kafka Connect

a. Running Kafka Connect

In this exercise, you will run Connect in distributed mode, and use the JDBC source Connector and File sink Connector. You will configure monitors using the REST API as well as the Control Center UI. You will use Control Center to monitor the connectors as well.

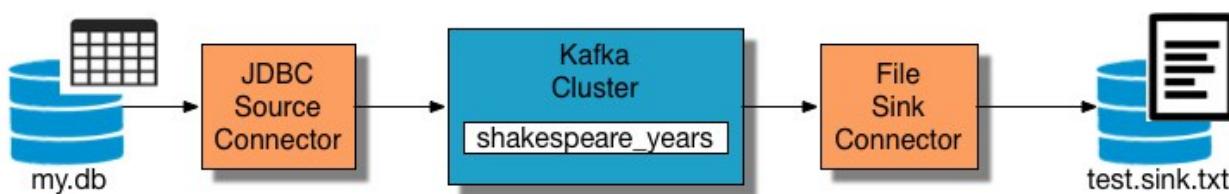
Prerequisites

1. Please make sure you have prepared your environment by following → [Preparing the Labs](#)
2. Start the Kafka cluster:

```
$ cd ~/confluent-admin  
$ docker-compose up -d
```

Connect Pipeline

In this section, you will run Connect in distributed mode with two Connectors: a JDBC source Connector and a file sink Connector. The JDBC source Connector writes the contents of a database table to a Kafka Topic. The file sink Connector reads data from the same Kafka Topic and writes those messages to a file. It will update when new rows are added to the database.



Prerequisites

1. Create the SQLite database **my.db** in the folder **~/confluent-admin/data**: (You may have to install SQLite - \$ sudo apt install sqlite3)

```
$ cd ~/confluent-admin  
$ sqlite3 data/my.db
```

2. Run the following statements in SQLite3> :

```
create table years(id INTEGER PRIMARY KEY AUTOINCREMENT, name  
VARCHAR(50), year INTEGER);  
insert into years(name,year) values('Hamlet',1600);  
insert into years(name,year) values('Julius Caesar',1599);  
insert into years(name,year) values('Macbeth',1605);  
insert into years(name,year) values('Merchant of Venice',1595);  
insert into years(name,year) values('Othello',1604);  
insert into years(name,year) values('Romeo and Juliette',1594);  
insert into years(name,year) values('Anthony and Cleopatra',1606);
```

3. Make sure the data is there:

```
sqlite> SELECT * FROM years;  
1|Hamlet|1600  
2|Julius Caesar|1599  
3|Macbeth|1605  
4|Merchant of Venice|1595  
5|Othello|1604  
6|Romeo and Juliette|1594  
7|Anthony and Cleopatra|1606
```

4. Press **Ctrl+D** to exit SQLite3.
5. Create a new Topic called **shakespeare_years** with one Partition and one replica.

```
$ kafka-topics \  
--bootstrap-server kafka-1:19092 \  
--create \  
--topic shakespeare_years \  
--partitions 1 \  
--replication-factor 1
```

Install the Kafka Connect JDBC Connector

We use the Kafka Connect JDBC connector in this exercise so we need to install it on the worker.

1. Install the connector:

```
$ docker-compose exec -u root connect confluent-hub install  
confluentinc/kafka-connect-jdbc:10.0.0
```

The component can be installed in any of the following Confluent Platform installations:

1. / (installed rpm/deb package)
2. / (where this tool is installed)

Choose one of these to continue the installation (1-2):

2. Type **1** and press **Enter**.

3. At the prompt, type **N** and press **Enter**.

```
Do you want to install this into /usr/share/confluent-hub-components?  
(yN)
```

4. At the prompt, type **/usr/share/java/kafka** and press **Enter**.

```
Specify installation directory:
```

5. At the prompt, type **y** and press **Enter**.

```
Component's license:
```

Confluent Community License

<https://www.confluent.io/confluent-community-license>

I agree to the software license agreement (yN)

6. At the prompt, type **y** and press **Enter**.

```
Downloading component Kafka Connect JDBC 10.0.0, provided by
Confluent, Inc. from Confluent Hub and installing into
/usr/share/java/kafka
Detected Worker's configs:
 1. Standard: /etc/kafka/connect-distributed.properties
 2. Standard: /etc/kafka/connect-standalone.properties
 3. Standard: /etc/schema-registry/connect-avro-
distributed.properties
 4. Standard: /etc/schema-registry/connect-avro-
standalone.properties
 5. Used by Connect process with PID : /etc/kafka-connect/kafka-
connect.properties
Do you want to update all detected configs? (yN)
```

The installation completes.

```
Adding installation directory to plugin path in the following files:
/etc/kafka/connect-distributed.properties
/etc/kafka/connect-standalone.properties
/etc/schema-registry/connect-avro-distributed.properties
/etc/schema-registry/connect-avro-standalone.properties
/etc/kafka-connect/kafka-connect.properties
```

Completed

7. To complete the installation, we need to restart the **connect** container:

```
$ docker-compose restart connect
```

8. Verify that the Connect Worker successfully restarted prior to continuing to the next step:

```
$ docker-compose logs connect | grep -i "INFO .* Finished starting
connectors and tasks"
connect | [2020-10-15 21:40:34,312] INFO [Worker clientId=connect-1,
groupId=connect] Finished starting connectors and tasks
(org.apache.kafka.connect.runtime.distributed.DistributedHerder:1236)
```



Repeat this command until the **Finished starting connectors and tasks** message appears.

Configuring the Source Connector

1. Add a new source connector to read data from the database **my.db** and write to the Kafka Topic **shakespeare_years** by using the Connect REST API:
 - a. Open Control Center at <http://localhost:9021>:
 - b. In the left sidebar, select **Cluster 1**
 - c. In the sidebar click **Connect**
 - d. In the view **All Connect Clusters** select the (only available) entry **connect-default**
 - e. Click **Add connector**
 - f. In the **Browse** overview select the **JdbcSourceConnector Source** tile
 - g. Configure the JDBC Source Connector
 - i. In the **Name** text box, type **JDBC-Source-Connector**
 - ii. In the **Database** section, in the **JDBC URL** text box, type
jdbc:sqlite:/data/my.db
 - iii. In the **Database** section, in the **Table Whitelist** dropdown, enter **years**
 - iv. Under section **Mode**, in the **Table Loading Mode** text box, type **incrementing**
 - v. Under section **Mode**, in the **Incrementing Column Name** text box, type **id**
 - vi. Under section **Connector**, in the **Topic Prefix** text box, type **shakespeare_**
 - h. Click **Continue**

CONFLUENT

CONTROL CENTER CLUSTER

Cluster 1

Overview

Brokers

Topics

Connect

ksqlDB

Consumers

Replicators

Cluster settings

01 SETUP CONNECTION

```
{
  "name": "JDBC-Source-Connector",
  "connector.class": "io.confluent.connect.jdbc.JdbcSourceConnector",
  "connection.url": "jdbc:sqlite:/data/my.db",
  "table.whitelist": [
    "years"
  ],
  "mode": "incrementing",
  "incrementing.column.name": "id",
  "topic.prefix": "shakespeare_"
}
```

02 TEST AND VERIFY

Launch **Back** [Download connector config file](#)

i. Click **Launch**

j. Verify you see the new connector **JDBC-Source-Connector** running

CONFLUENT

Enterprise trial ends in 29 days

CONNECT CLUSTERS > CONNECT-DEFAULT >

Connectors

Total	Running	Degraded	Failed	Paused
1	1	0	0	0

Search connectors Filter by category + Add connector [Upload connector config file](#)

Status	Name	Category	Plugin name	Topics	Number of tasks
Running	JDBC-Source-Connector	Source	JdbcSourceConnector	--	1

k. **Optional:** As an alternative to using Control Center, you could instead add the source connector via command line using Kafka Connect's **REST API**:

```
$ curl -s -X POST \
-H "Content-Type: application/json" \
--data '{
    "name": "JDBC-Source-Connector",
    "config": {
        "connector.class":
"io.confluent.connect.jdbc.JdbcSourceConnector",
        "connection.url": "jdbc:sqlite:/data/my.db",
        "table.whitelist": "years",
        "mode": "incrementing",
        "incrementing.column.name": "id",
        "table.types": "TABLE",
        "topic.prefix": "shakespeare_"
    }
}' http://connect:8083/connectors
```

2. Launch another terminal and start the console Consumer for Topic **shakespeare_years**:

```
$ kafka-console-consumer \
--bootstrap-server kafka-1:19092, kafka-2:29092 \
--property schema.registry.url=http://schema-registry:8081 \
--from-beginning \
--topic shakespeare_years

{"id":1,"name":{"string":"Hamlet"},"year":{"long":1600}}
{"id":2,"name":{"string":"Julius Caesar"},"year":{"long":1599}}
{"id":3,"name":{"string":"Macbeth"},"year":{"long":1605}}
{"id":4,"name":{"string":"Merchant of Venice"},"year":{"long":1595}}
{"id":5,"name":{"string":"Othello"},"year":{"long":1604}}
{"id":6,"name":{"string":"Romeo and Juliette"},"year":{"long":1594}}
{"id":7,"name":{"string":"Anthony and Cleopatra"},"year":{"long":1606}}
...
```

Leave the Consumer running until instructed to terminate it. See what messages have been and will be produced.

Configuring the Sink Connector

1. Open a second terminal window and navigate to the **confluent-admin** folder:

```
$ cd ~/confluent-admin
```

2. Update permissions for the **data** directory to allow access by the **connect** container:

```
$ chmod 777 data
```

3. Add a new sink connector to read data from the Kafka Topic **shakespeare_years** and write to the file **data/test.sink.txt** on the Connect worker system:

```
$ curl -s -X POST \
  -H "Content-Type: application/json" \
  --data '{
    "name": "File-Sink-Connector",
    "config": {
      "topics": "shakespeare_years",
      "connector.class":
"org.apache.kafka.connect.file.FileStreamSinkConnector",
      "value.converter":
"io.confluent.connect.avro.AvroConverter",
      "value.converter.schema.registry.url":
"http://schema-registry:8081",
      "file": "/data/test.sink.txt"
    }
  }' http://connect:8083/connectors
```

4. Use Control Center to verify the creation of the connector.

Status	Name	Category	Plugin name	Topics	Number of tasks
Running	JDBC-Source-Connector	Source	JdbcSourceConnector	--	1
Running	File-Sink-Connector	Sink	FileStreamSinkConnector	shakespeare_years	1

5. Verify that a new file called **test.sink.txt** has been created in the folder **~/confluent-admin/data**. View this file to confirm that the sink connector worked:

```
$ cat data/test.sink.txt
```

```
Struct{id=1,name=Hamlet,year=1600}  
Struct{id=2,name=Julius Caesar,year=1599}  
Struct{id=3,name=Macbeth,year=1605}  
Struct{id=4,name=Merchant of Venice,year=1596}  
Struct{id=5,name=Othello,year=1604}  
Struct{id=6,name=Romeo and Juliet,year=1594}  
Struct{id=7,name=Antony and Cleopatra,year=1606}
```



This file was created by the **File Sink Connector** in folder `/data/` on the `connect` container. Since this container folder is mapped to the folder `~/confluent-admin/data` of your host system, we can see it there too.

6. Insert a few new rows into the table `years`:

- From your host, run SQLite3:

```
$ sqlite3 data/my.db
```

- Insert two records:

```
sqlite> INSERT INTO years(name,year) VALUES('Tempest',1611);  
        INSERT INTO years(name,year) VALUES('King Lear',1605);
```

- In the window where the Consumer is still running observe that two new records have been output:

```
{"id":8,"name":{"string":"Tempest"},"year":{"long":1611}}  
{"id":9,"name":{"string":"King Lear"},"year":{"long":1605}}
```

- Quit SQLite3 by pressing **Ctrl+D**.

7. View the sink file, `test.sink.txt` again and notice that the new log lines are added to it:

```
$ cat data/test.sink.txt
```

```
Struct{id=1,name=Hamlet,year=1600}  
Struct{id=2,name=Julius Caesar,year=1599}  
Struct{id=3,name=Macbeth,year=1605}  
Struct{id=4,name=Merchant of Venice,year=1596}  
Struct{id=5,name=Othello,year=1604}  
Struct{id=6,name=Romeo and Juliet,year=1594}  
Struct{id=7,name=Antony and Cleopatra,year=1606}  
Struct{id=8,name=Tempest,year=1611}  
Struct{id=9,name=King Lear,year=1605}
```

8. In Control Center, observe the Consumer Group performance for Kafka Connect. You could use this to ensure that Kafka Connect is performing well in your cluster, just like other production traffic.

- a. In Control Center, make sure that you have selected **Cluster 1** and then click on **Consumers**

Consumer group ID	Messages behind	Number of consumers	Number of topics
_confluent-controlcenter-5-5-0-1-command	0	1	1
connect-File-Sink-Connector	0	1	1
_confluent-controlcenter-5-5-0-1	7,516	4	15

- b. In the list of Consumer Groups select **connect-File-Sink-Connector**. You will then see the **Consumer lag** view

ALL CONSUMER GROUPS >

connect-File-Sink-Connector

Consumer lag **Consumption**

0 Total Messages behind +0 messages
5 second interval

shakespeare_years
Max lag / consumer: 0 messages

Consumer ID	Topic Partition	Partition	Lag Messages behind	Current offset	Error count
connector-consumer-File-Sink-Connector-0...	Shakespeare_years_0	0	0	16	0

9. In the terminal window where the consumer is running, press **Ctrl+C** to terminate the process.

Cleanup

1. Execute the following command to completely cleanup your environment. If you wish to preserve the data in the cluster so that you can continue to experiment, omit the **-v** option.

```
$ docker-compose down -v
```



STOP HERE. THIS IS THE END OF THE EXERCISE.

Appendix A: Reassigning Partitions in a Topic - Alternate Method

In the **Kafka Administrative Tools** exercise, one of the sections uses the Confluent Auto Data Balancer to rebalance the cluster and to reassign Partitions for the Topic **moving**. If you prefer to use the Apache open source tool **kafka-reassign-partitions** instead, you may follow the instructions below.



This appendix assumes that the Topic **moving** exists and has been populated as described in the **Kafka Administrative Tools** exercise. If this is not the case then please first follow the instructions → [Rebalancing the Cluster](#).

1. Change to the **data** folder:

```
$ cd ~/confluent-admin/data
```

2. In that folder create a file **topics-to-move.json** as follows:

```
$ echo '{"topics": [{"topic": "moving"}],"version":1}' > topics-to-move.json
```

3. Generate the reassignment plan.

```
$ kafka-reassign-partitions \
--bootstrap-server kafka-1:19092 \
--topics-to-move-json-file topics-to-move.json \
--broker-list "101,102,103" \
--generate > reassignment.json
```

The content of the file **reassignment.json** should look like this:

```
$ cat reassignment.json
```

Current partition replica assignment

```
{"version":1,"partitions":[{"topic":"moving","partition":2,"replicas":[101,102],"log_dirs":["any","any"]},{ "topic":"moving","partition":3,"replicas":[102,101],"log_dirs":["any","any"]},{ "topic":"moving","partition":0,"replicas":[101,102],"log_dirs":["any","any"]},{ "topic":"moving","partition":4,"replicas":[101,102],"log_dirs":["any","any"]},{ "topic":"moving","partition":5,"replicas":[102,101],"log_dirs":["any","any"]},{ "topic":"moving","partition":1,"replicas":[102,101],"log_dirs":["any","any"]}]}
```

Proposed partition reassignment configuration

```
{"version":1,"partitions":[{"topic":"moving","partition":2,"replicas":[101,102],"log_dirs":["any","any"]},{ "topic":"moving","partition":4,"replicas":[103,102],"log_dirs":["any","any"]},{ "topic":"moving","partition":1,"replicas":[103,101],"log_dirs":["any","any"]},{ "topic":"moving","partition":0,"replicas":[102,103],"log_dirs":["any","any"]},{ "topic":"moving","partition":3,"replicas":[102,101],"log_dirs":["any","any"]},{ "topic":"moving","partition":5,"replicas":[101,103],"log_dirs":["any","any"]}]}
```

4. Open the file **~/confluent-admin/data/reassignment.json** and edit it so it only includes the "Proposed partition reassignment configuration" in JSON (i.e., just the last line).

The resulting file should look like this:

```
$ cat reassignment.json
```

```
{"version":1,"partitions":[{"topic":"moving","partition":2,"replicas":[101,102],"log_dirs":["any","any"]},{ "topic":"moving","partition":4,"replicas":[103,102],"log_dirs":["any","any"]},{ "topic":"moving","partition":1,"replicas":[103,101],"log_dirs":["any","any"]},{ "topic":"moving","partition":0,"replicas":[102,103],"log_dirs":["any","any"]},{ "topic":"moving","partition":3,"replicas":[102,101],"log_dirs":["any","any"]},{ "topic":"moving","partition":5,"replicas":[101,103],"log_dirs":["any","any"]}]}
```

5. Observe the current throttling limits configured:

```
$ kafka-configs \
  --describe \
  --bootstrap-server kafka-1:19092 \
  --entity-type brokers
Dynamic configs for broker 101 are:
Dynamic configs for broker 102 are:
Dynamic configs for broker 103 are:
Default configs for brokers in the cluster are:
```



kafka-configs by default only reports dynamic configuration settings. Since none have been set at this time, the dynamic configs list for each broker should be empty.

6. Execute the reassignment plan with throttling set to 1MBps.

```
$ kafka-reassign-partitions \
  --bootstrap-server kafka-1:19092 \
  --reassignment-json-file reassignment.json \
  --execute \
  --throttle 1000000
```

Current partition replica assignment

```
{"version":1,"partitions":[{"topic":"moving","partition":2,"replicas":[101,102],"log_dirs":["any","any"]},{ "topic":"moving","partition":3,"replicas":[102,101],"log_dirs":["any","any"]},{ "topic":"moving","partition":0,"replicas":[101,102],"log_dirs":["any","any"]},{ "topic":"moving","partition":4,"replicas":[101,102],"log_dirs":["any","any"]},{ "topic":"moving","partition":5,"replicas":[102,101],"log_dirs":["any","any"]},{ "topic":"moving","partition":1,"replicas":[102,101],"log_dirs":["any","any"]}]}
```

Save this to use as the --reassignment-json-file option during rollback

Warning: You must run Verify periodically, until the reassignment completes, to ensure the throttle is removed. You can also alter the throttle by rerunning the Execute command passing a new value.

The inter-broker throttle limit was set to 1000000 B/s

Successfully started reassignment of partitions.

7. Observe the new throttling limits configured.

```

$ kafka-configs \
  --describe \
  --bootstrap-server kafka-1:19092 \
  --entity-type brokers
Dynamic configs for broker 101 are:
  leader.replication.throttled.rate=1000000 sensitive=false
  synonyms={DYNAMIC_BROKER_CONFIG:leader.replication.throttled.rate=100000}
  follower.replication.throttled.rate=1000000 sensitive=false
  synonyms={DYNAMIC_BROKER_CONFIG:follower.replication.throttled.rate=1000000}
Dynamic configs for broker 102 are:
  leader.replication.throttled.rate=1000000 sensitive=false
  synonyms={DYNAMIC_BROKER_CONFIG:leader.replication.throttled.rate=100000}
  follower.replication.throttled.rate=1000000 sensitive=false
  synonyms={DYNAMIC_BROKER_CONFIG:follower.replication.throttled.rate=1000000}
Dynamic configs for broker 103 are:
  leader.replication.throttled.rate=1000000 sensitive=false
  synonyms={DYNAMIC_BROKER_CONFIG:leader.replication.throttled.rate=100000}
  follower.replication.throttled.rate=1000000 sensitive=false
  synonyms={DYNAMIC_BROKER_CONFIG:follower.replication.throttled.rate=1000000}
Default configs for brokers in the cluster are:

```

8. Monitor the progress of the reassignment.

- Run the **kafka-reassign-partitions** command with the **--verify** option and note that some reassessments are "still in progress".

```

$ kafka-reassign-partitions \
  --bootstrap-server kafka-1:19092 \
  --reassignment-json-file reassignment.json \
  --verify

```

Status of partition reassignment:

Reassignment of partition moving-2 completed successfully
 Reassignment of partition moving-3 completed successfully
 Reassignment of partition moving-0 is still in progress
 Reassignment of partition moving-4 is still in progress
 Reassignment of partition moving-5 is still in progress
 Reassignment of partition moving-1 is still in progress

- Run the **kafka-topics** command with the **--describe** option and notice the listed replicas for the Partitions that are still in progress for reassignment.

```

$ kafka-topics \
    --bootstrap-server kafka-1:19092 \
    --describe \
    --topic moving

Topic:moving      PartitionCount:6      ReplicationFactor:3
Configs:leader.replication.throttled.replicas=0:101,0:102,4:101,4:
102,5:102,5:101,1:102,1:101,follower.replication.throttled.replica
s=0:103,4:103,5:103,1:103
    Topic: moving      Partition: 0      Leader: 101 Replicas:
102,103,101      Isr: 101,102
    Topic: moving      Partition: 1      Leader: 102 Replicas:
103,101,102      Isr: 102,101
    Topic: moving      Partition: 2      Leader: 101 Replicas: 101,102
Isr: 101,102
    Topic: moving      Partition: 3      Leader: 102 Replicas: 102,101
Isr: 102,101
    Topic: moving      Partition: 4      Leader: 101 Replicas:
103,102,101      Isr: 101,102
    Topic: moving      Partition: 5      Leader: 102 Replicas:
101,103,102      Isr: 102,101

```

9. Increase the throttle limit configuration to 1GBps by rerunning the **kafka-reassign-partitions** command with the new throttle limit.

```

$ kafka-reassign-partitions \
    --bootstrap-server kafka-1:19092 \
    --reassignment-json-file reassignment.json \
    --execute \
    --throttle 1000000000 \
    --additional
Current partition replica assignment

{"version":1,"partitions":[{"topic":"moving","partition":0,"replicas":[103,101,102],"log_dirs":["any","any","any"]},{ "topic":"moving","partition":1,"replicas":[101,102],"log_dirs":["any","any"]},{ "topic":"mo
ving","partition":2,"replicas":[102,103,101],"log_dirs":["any","any","any"]},{ "topic":"moving","partition":3,"replicas":[103,102,101],"log
_dirs":["any","any","any"]},{ "topic":"moving","partition":4,"replicas":[101,103,102],"log_dirs":["any","any","any"]},{ "topic":"moving","pa
rtition":5,"replicas":[102,101],"log_dirs":["any","any"]}]}
```

Save this to use as the --reassignment-json-file option during rollback

Warning: You must run --verify periodically, until the reassignment completes, to ensure the throttle is removed.

The inter-broker throttle limit was set to 1000000000 B/s

10. Run the **kafka-reassign-partitions** command with the **--verify** option again. Note the "completed successfully" output and "Throttle was removed".

```
$ kafka-reassign-partitions \
  --bootstrap-server kafka-1:19092 \
  --reassignment-json-file reassignment.json \
  --verify
Status of partition reassignment:
Reassignment of partition moving-0 is complete.
Reassignment of partition moving-1 is complete.
Reassignment of partition moving-2 is complete.
Reassignment of partition moving-3 is complete.
Reassignment of partition moving-4 is complete.
Reassignment of partition moving-5 is complete.

Clearing broker-level throttles on brokers 101,102,103
Clearing topic-level throttles on topic moving
```

11. Confirm that the throttle limit is now removed.

```
$ kafka-configs \
  --describe \
  --bootstrap-server kafka-1:19092 \
  --entity-type brokers
Dynamic configs for broker 101 are:
Dynamic configs for broker 102 are:
Dynamic configs for broker 103 are:
Default configs for brokers in the cluster are:
```

12. Return to (the end of) the exercise **Rebalancing the Cluster** (\rightarrow [Simulate a Completely Failed Broker](#)).



STOP HERE. THIS IS THE END OF THE EXERCISE.

Appendix B: Running Labs in Docker for Desktop

If you have installed Docker for Desktop on your Mac or Windows 10 Pro machine you are able to complete the course by building and running your applications from the command line.

- Increase the memory available to Docker Desktop to a minimum of 8 GiB. See the advanced settings for [Docker Desktop for Mac](#), and [Docker Desktop for Windows](#).
- Follow the instructions at → [Preparing the Labs](#) to `git clone` the source code, and launch the cluster containers with `docker-compose`.
- In each exercise, open a bash shell in the **tools container** to run commands against the cluster. All the command line instructions will work from the **tools** container. This container has been preconfigured with all of the tools you use in the exercises, e.g. **kafka-topics**, **confluent-rebalancer**, and **kafka-configs**. You can think of this container as a "bastion" inside of your distributed architecture.

```
$ docker-compose exec tools bash  
root@tools:#
```



The `~/confluent-admin/data` folder is mapped to `/apps/data` in the **tools** container, so any directions that use `~/confluent-admin/data` should be executed from `/apps/data` inside the **tools** container.

- Anywhere you are instructed to open additional terminal windows, you can open additional bash shells on the **tools** container with the same command as above on your host machine.
- Any subsequent `docker` or `docker-compose` instructions should be run on your host machine.